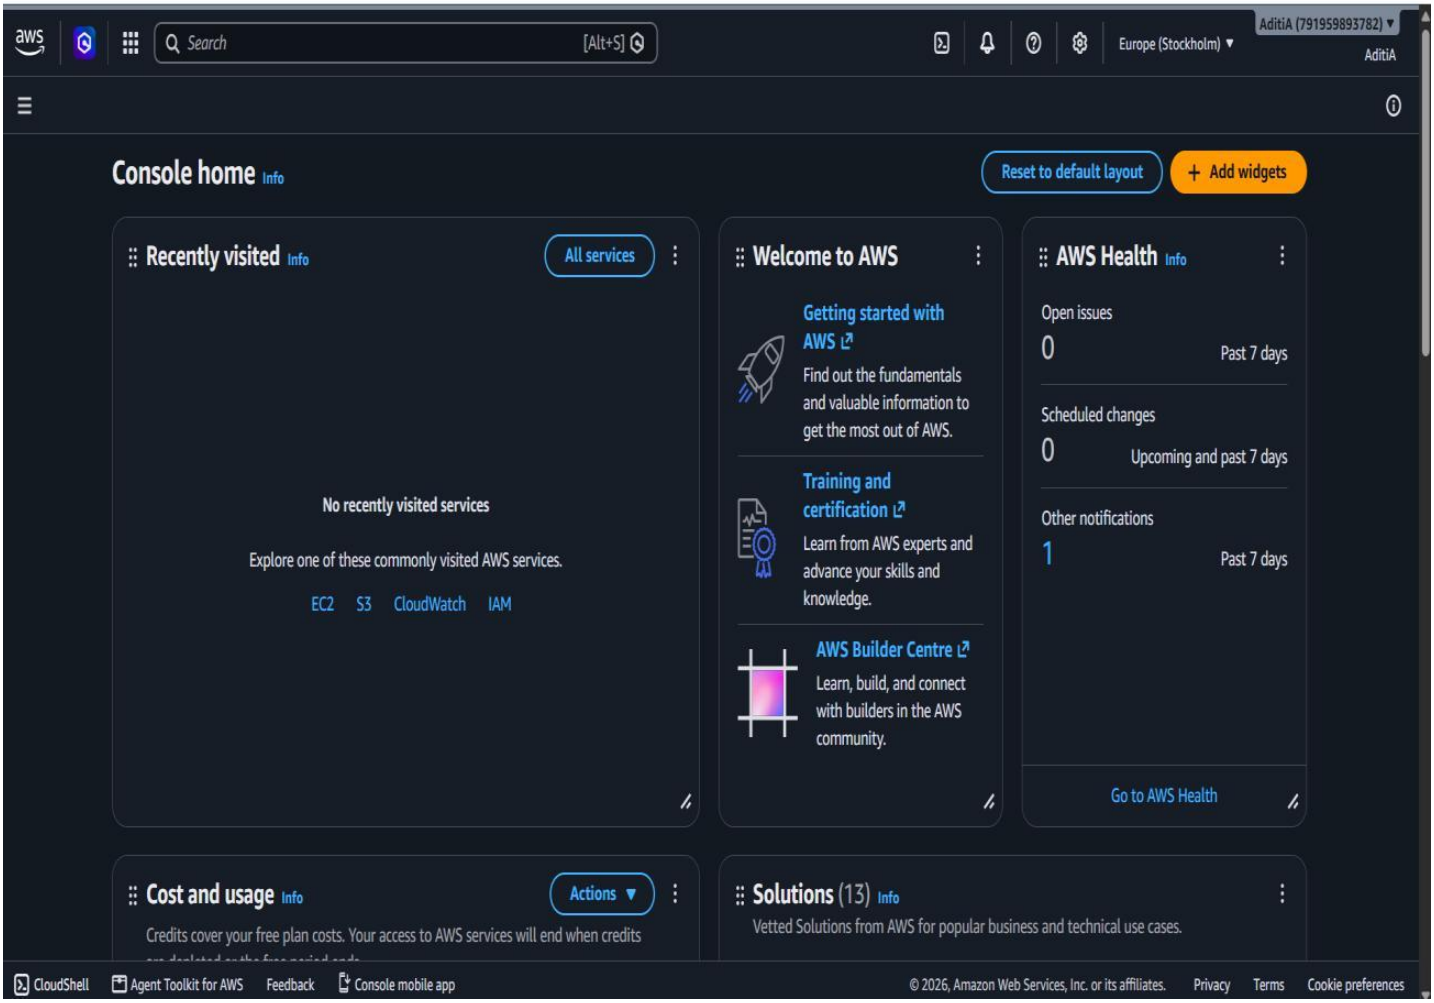
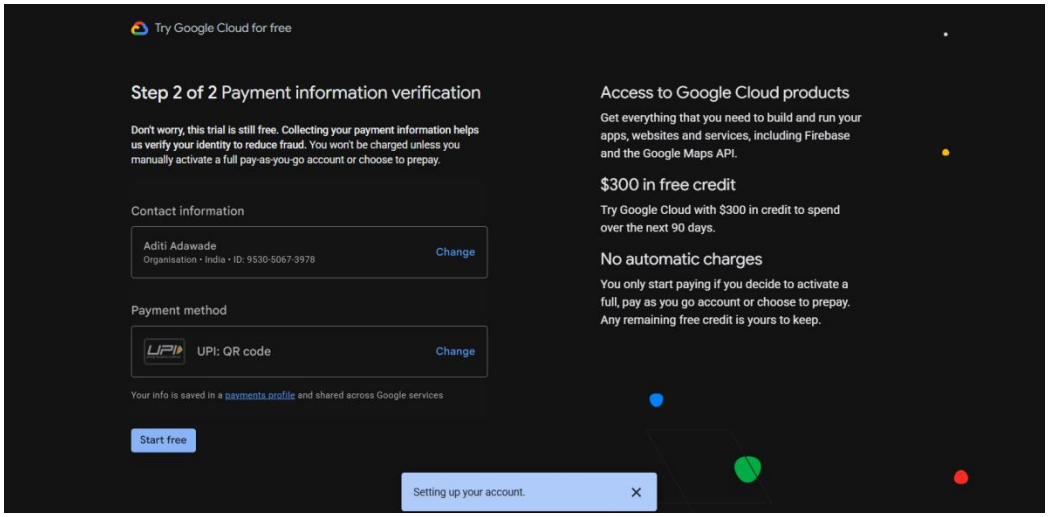
AWS Console home dashboard

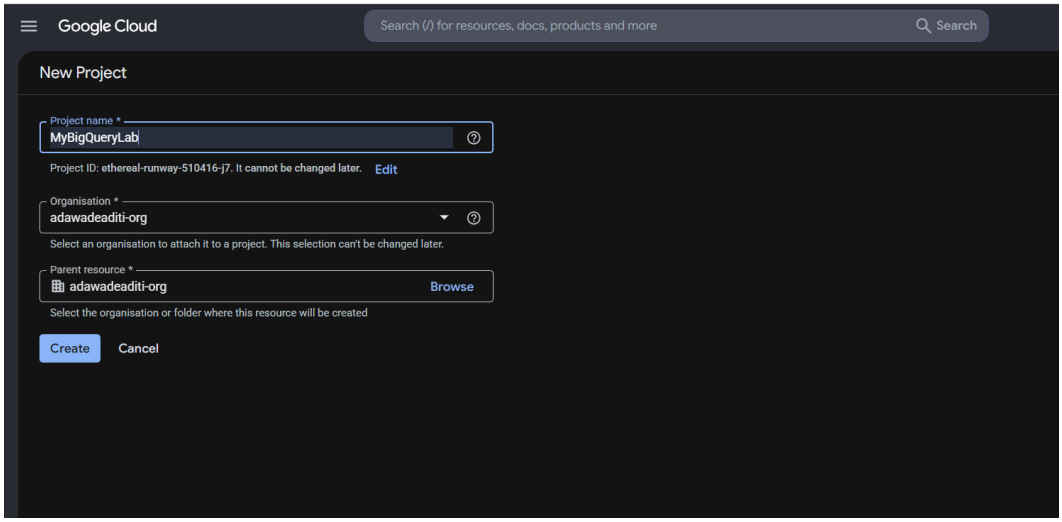


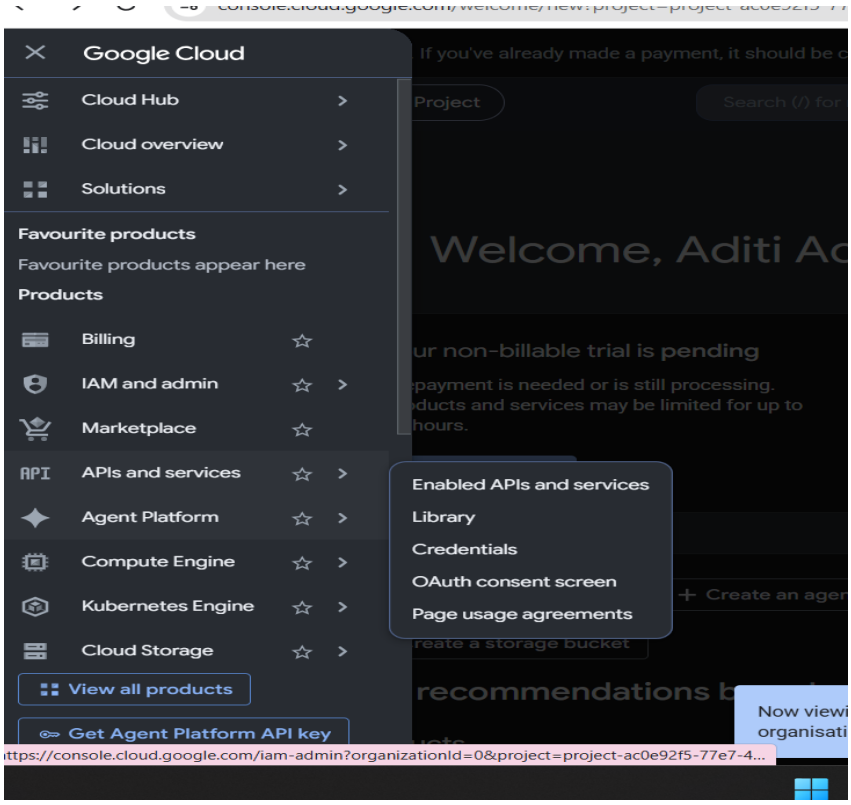
CCD Practical no. 5

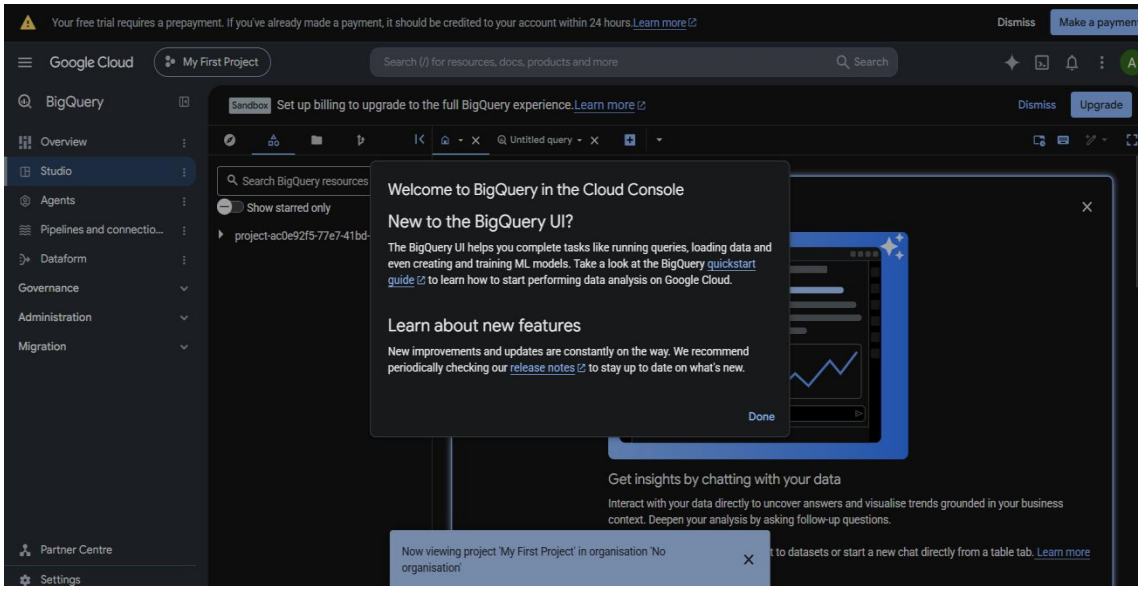
Step 1:Create an account on google cloud platform



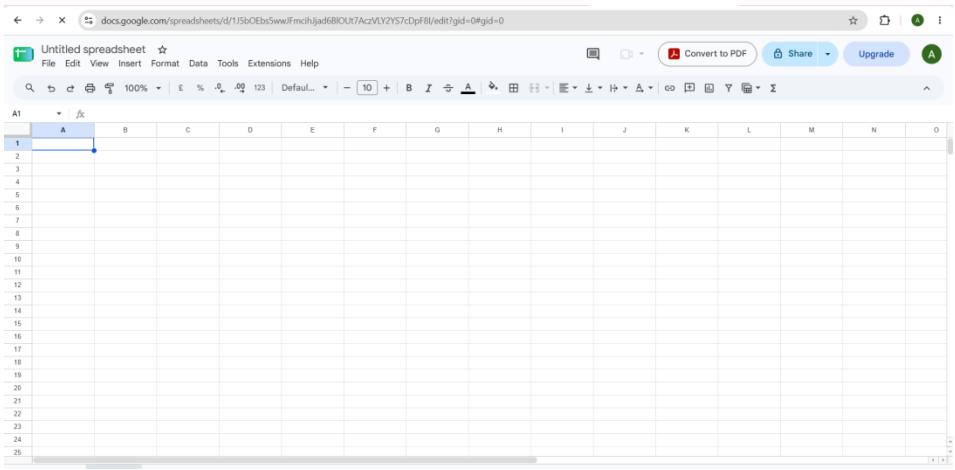
Prat 2: Create a project and access BigQuery




Google Cloud
Cloud Hub
Cloud overview
Solutions
Favourite products
Favourite products appear here
Products
Billing
IAM and admin
Marketplace
APIs and services
Agent Platform
Compute Engine
Kubernetes Engine
Cloud Storage
View all products
Get Agent Platform API key
Enabled APIs and services
Library
Credentials
OAuth consent screen
Page usage agreements
Welcome, Aditi
https://console.cloud.google.com/iam-admin?organizationId=0&project=project-ac0e92f5-77e7-4...


Your free trial requires a prepayment. If you've already made a payment, it should be credited to your account within 24 hours. Learn more
Dismiss
Make a payment
Google Cloud
My First Project
Search (/) for resources, docs, products and more
Search
BigQuery
Sandbox
Set up billing to upgrade to the full BigQuery experience. Learn more
Dismiss
Upgrade
Overview
Studio
Agents
Pipelines and connectio...
Dataform
Governance
Administration
Migration
Partner Centre
Settings
Search BigQuery resources
Show starred only
project-ac0e92f5-77e7-41bd-
Untitled query
Welcome to BigQuery in the Cloud Console
New to the BigQuery UI?
The BigQuery UI helps you complete tasks like running queries, loading data and even creating and training ML models. Take a look at the BigQuery quickstart guide to learn how to start performing data analysis on Google Cloud.
Learn about new features
New improvements and updates are constantly on the way. We recommend periodically checking our release notes to stay up to date on what's new.
Done
Get insights by chatting with your data
Interact with your data directly to uncover answers and visualise trends grounded in your business context. Deepen your analysis by asking follow-up questions.
Now viewing project 'My First Project' in organisation 'No organisation'
to datasets or start a new chat directly from a table tab. Learn more

Part 3: Query data directly from google sheets

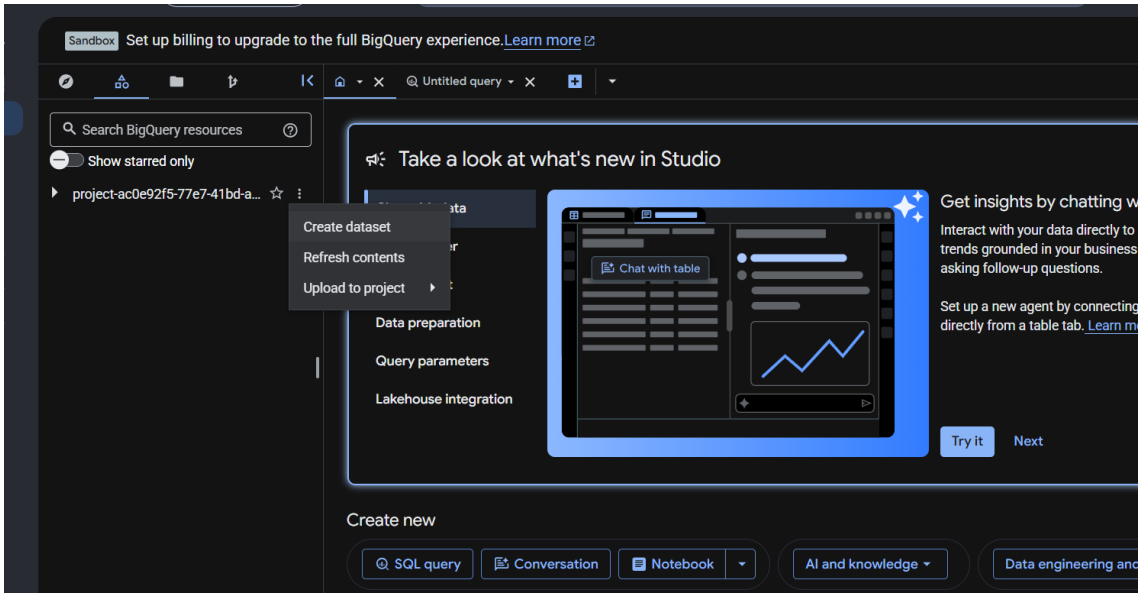


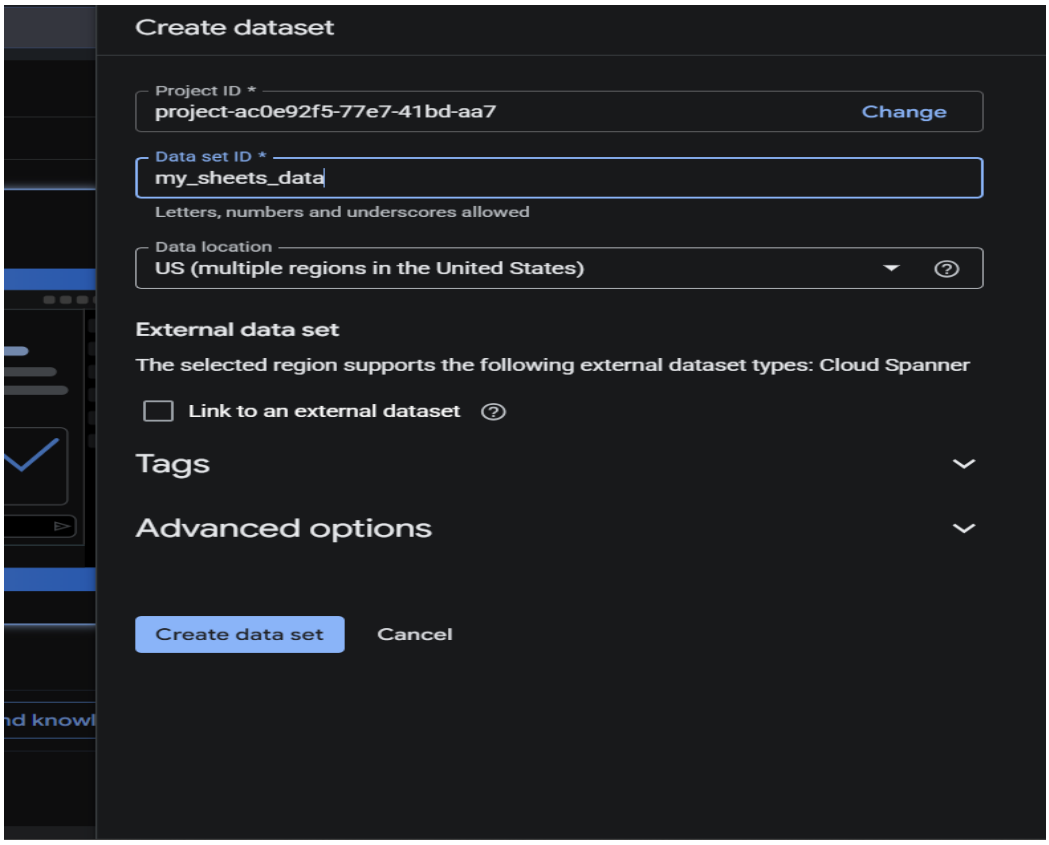
Product_Data

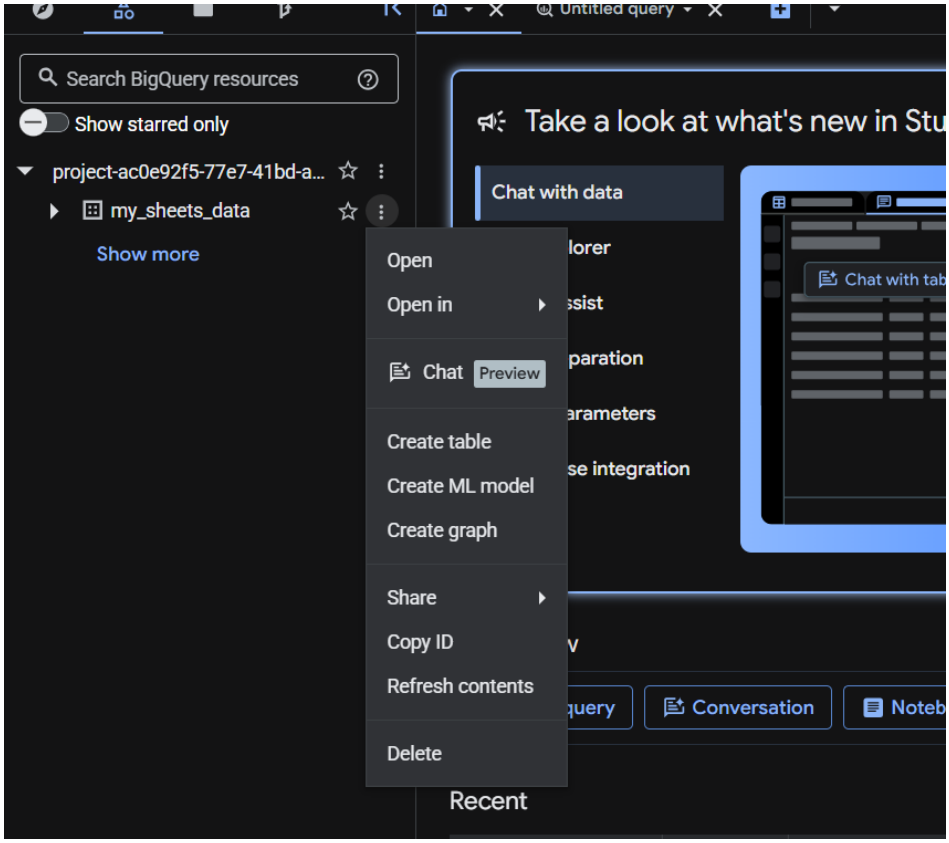
File Edit View Insert Format Data Tools Extensions Help

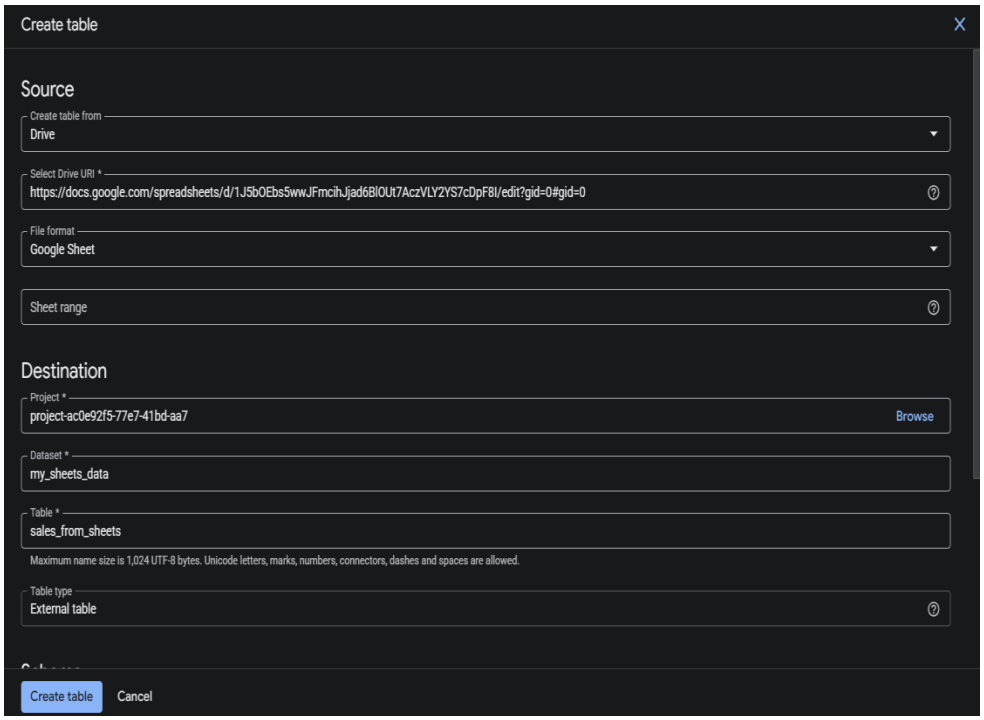
I10

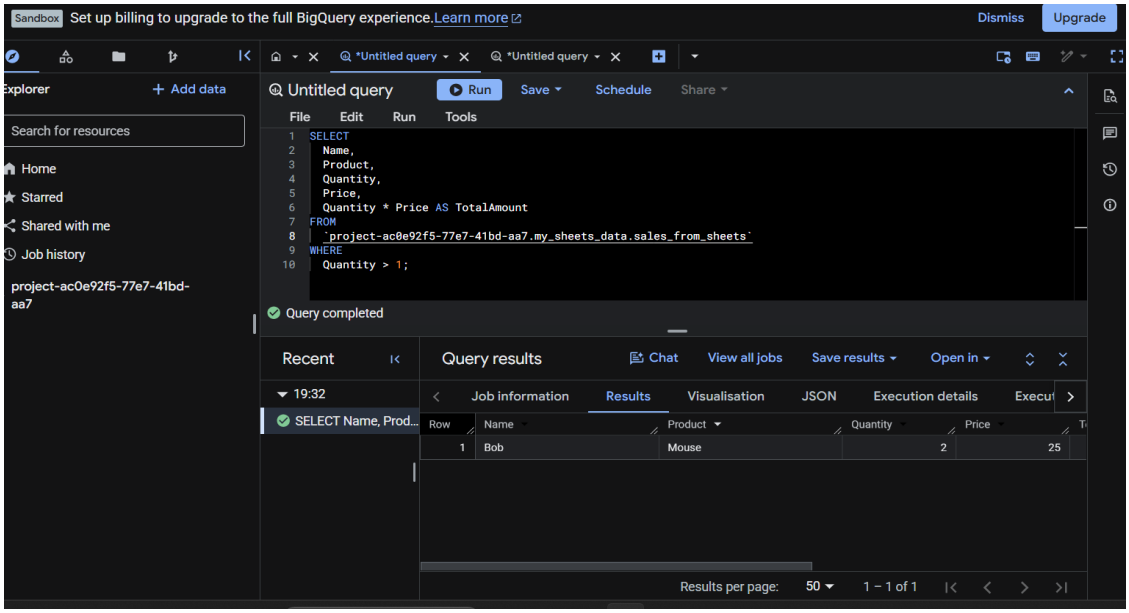
	A	B	C	D	E	F	G	H
1	Name	Product	Quantity	Price				
2	Alice	Laptop	1	1200				
3	Bob	Mouse	2	25				
4	Charlie	Keyboard	1	75				
5	Alice	Monitor	1	300				
6								
7								
8								
9								
10								
11								
12								
13								
14								
15								
16								


Sandbox
Set up billing to upgrade to the full BigQuery experience.Learn more
Untitled query
Search BigQuery resources
Show starred only
project-ac0e92f5-77e7-41bd-a...
Create dataset
Refresh contents
Upload to project
Take a look at what's new in Studio
Data preparation
Query parameters
Lakehouse integration
Chat with table
Get insights by chatting w
Interact with your data directly to
trends grounded in your business
asking follow-up questions.
Set up a new agent by connecting
directly from a table tab. Learn m
Try it
Next
Create new
SQL query
Conversation
Notebook
AI and knowledge
Data engineering anc

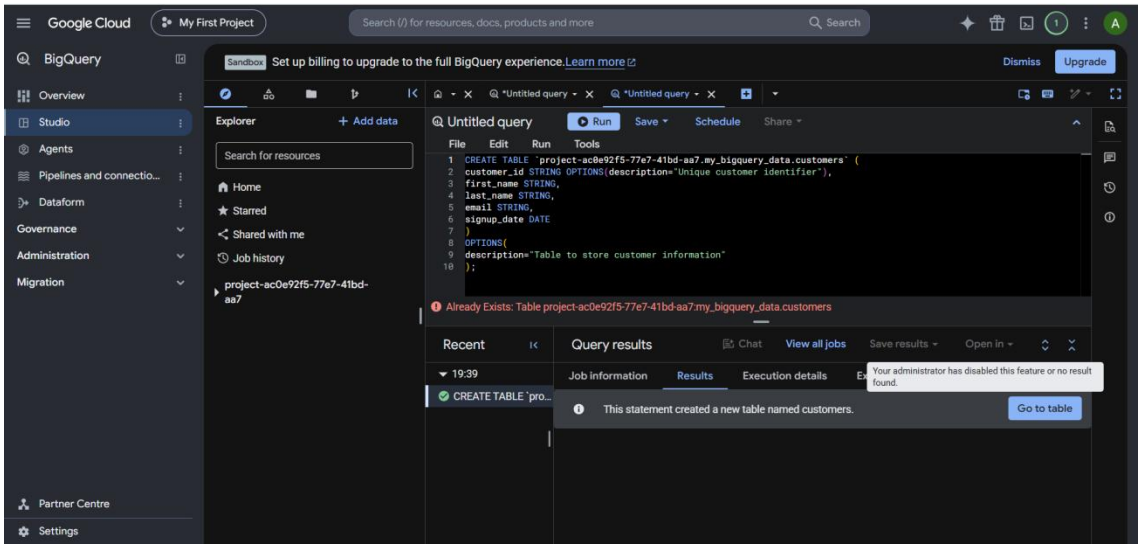

Create dataset
Project ID *
project-ac0e92f5-77e7-41bd-aa7
Change
Data set ID *
my_sheets_data
Letters, numbers and underscores allowed
Data location
US (multiple regions in the United States)
External data set
The selected region supports the following external dataset types: Cloud Spanner
Link to an external dataset
Tags
Advanced options
Create data set
Cancel

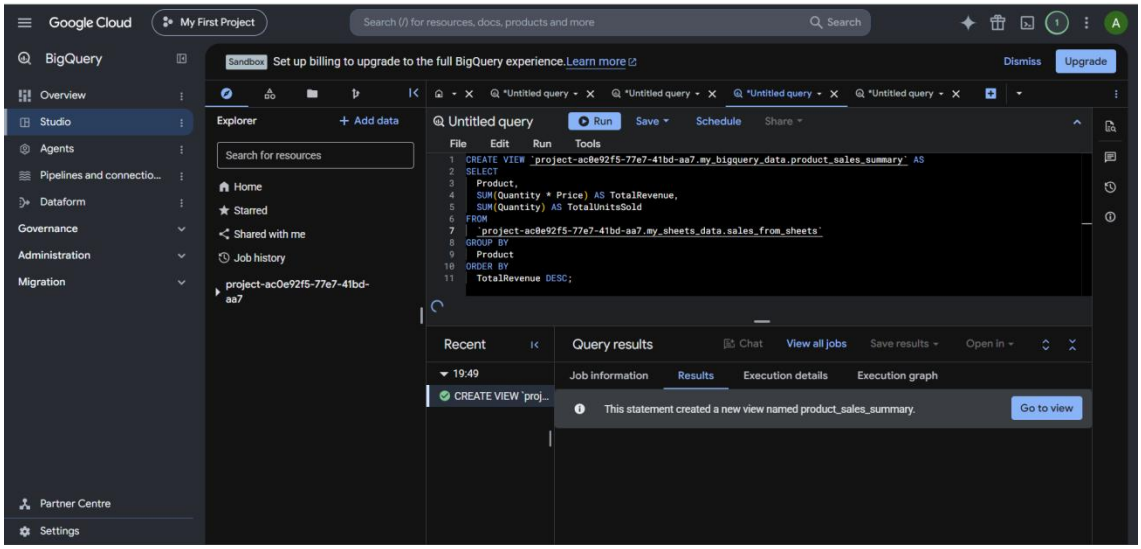

Untitled query
Search BigQuery resources
Show starred only
project-ac0e92f5-77e7-41bd-a...
my_sheets_data
Show more
Open
Open in
Chat
Preview
Create table
Create ML model
Create graph
Share
Copy ID
Refresh contents
Delete
Take a look at what's new in Stu
Chat with data
Chat with tab
Conversation
Recent

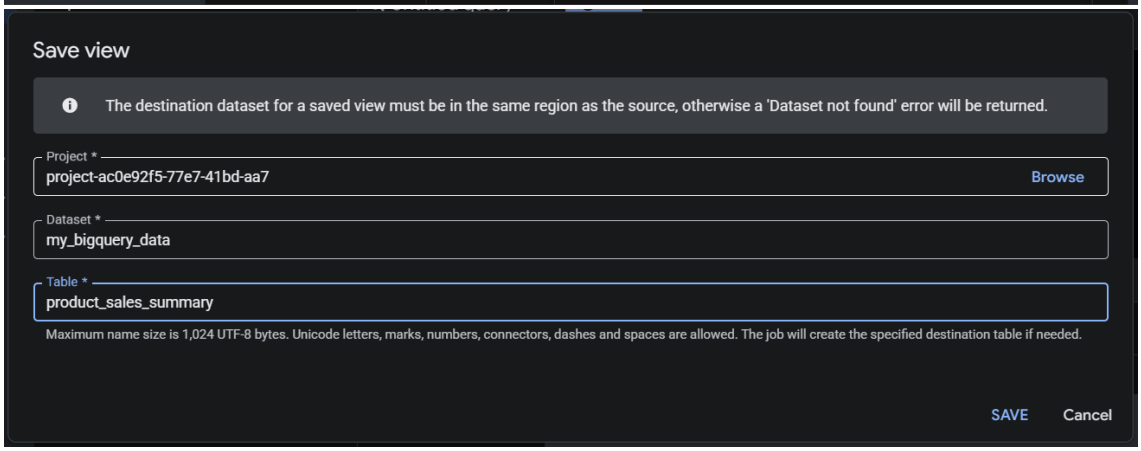

Create table
Source
Create table from
Drive
Select Drive URI *
https://docs.google.com/spreadsheets/d/1J5bOEbs5wwJFmcihJjad6BlOUt7AczVLY2YS7cDpF8I/edit?gid=0#gid=0
File format
Google Sheet
Sheet range
Destination
Project *
project-ac0e92f5-77e7-41bd-aa7
Browse
Dataset *
my_sheets_data
Table *
sales_from_sheets
Maximum name size is 1,024 UTF-8 bytes. Unicode letters, marks, numbers, connectors, dashes and spaces are allowed.
Table type
External table
Create table
Cancel

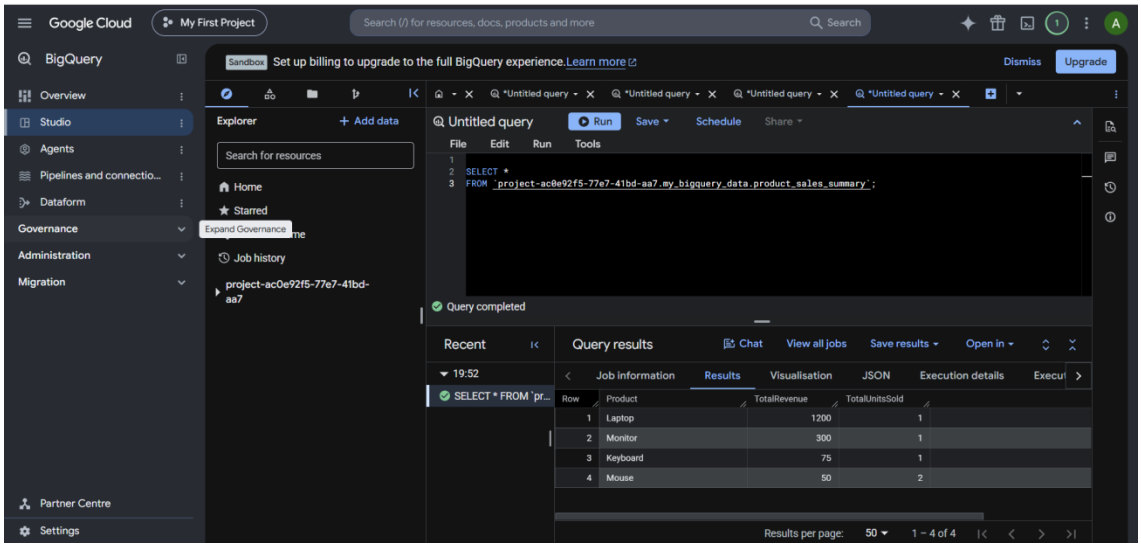


Part 4: Create Tables and views using BigQuery




Google Cloud
My First Project
Search (/) for resources, docs, products and more
BigQuery
Sandbox
Set up billing to upgrade to the full BigQuery experience. Learn more
Dismiss
Upgrade
Overview
Studio
Agents
Pipelines and connectio...
Dataform
Governance
Administration
Migration
Explorer
+ Add data
Search for resources
Home
Starred
Shared with me
Job history
project-ac0e92f5-77e7-41bd-aa7
Untitled query
Run
Save
Schedule
Share
File
Edit
Run
Tools
CREATE VIEW `project-ac0e92f5-77e7-41bd-aa7.my_bigquery_data.product_sales_summary` AS
SELECT
Product,
SUM(Quantity * Price) AS TotalRevenue,
SUM(Quantity) AS TotalUnitsSold
FROM
`project-ac0e92f5-77e7-41bd-aa7.my_sheets_data.sales_from_sheets`
GROUP BY
Product
ORDER BY
TotalRevenue DESC;
Recent
Query results
Chat
View all jobs
Save results
Open in
19:49
CREATE VIEW `proj...
Job information
Results
Execution details
Execution graph
This statement created a new view named product_sales_summary.
Go to view
Partner Centre
Settings

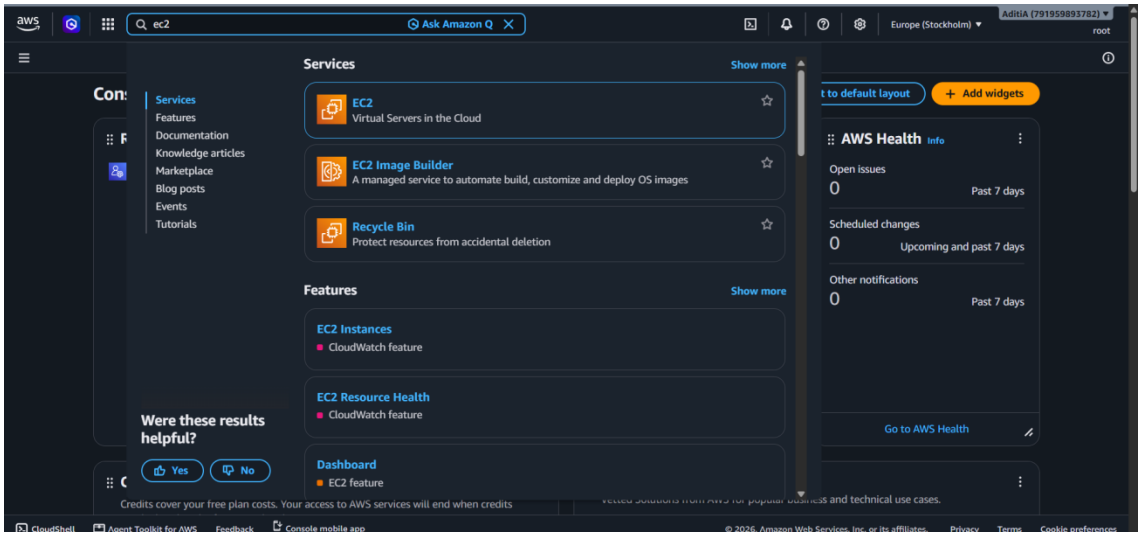

Save view
The destination dataset for a saved view must be in the same region as the source, otherwise a 'Dataset not found' error will be returned.
Project *
project-ac0e92f5-77e7-41bd-aa7
Browse
Dataset *
my_bigquery_data
Table *
product_sales_summary
Maximum name size is 1,024 UTF-8 bytes. Unicode letters, marks, numbers, connectors, dashes and spaces are allowed. The job will create the specified destination table if needed.
SAVE
Cancel


Google Cloud
My First Project
Search (/) for resources, docs, products and more
BigQuery
Sandbox
Set up billing to upgrade to the full BigQuery experience. Learn more
Dismiss
Upgrade
Overview
Studio
Agents
Pipelines and connectio...
Dataform
Governance
Expand Governance
Administration
Migration
Explorer
+ Add data
Search for resources
Home
Starred
Job history
project-ac0e92f5-77e7-41bd-aa7
Untitled query
Run
Save
Schedule
Share
File
Edit
Run
Tools
SELECT *
FROM `project-ac0e92f5-77e7-41bd-aa7.my_bigquery_data.product_sales_summary`;
Query completed
Recent
Query results
Chat
View all jobs
Save results
Open in
19:52
SELECT * FROM `pr...
Job information
Results
Visualisation
JSON
Execution details
Row
Product
TotalRevenue
TotalUnitsSold
1
Laptop
1200
1
2
Monitor
300
1
3
Keyboard
75
1
4
Mouse
50
2
Results per page:
50
1 – 4 of 4
Partner Centre
Settings

Practical no. 6

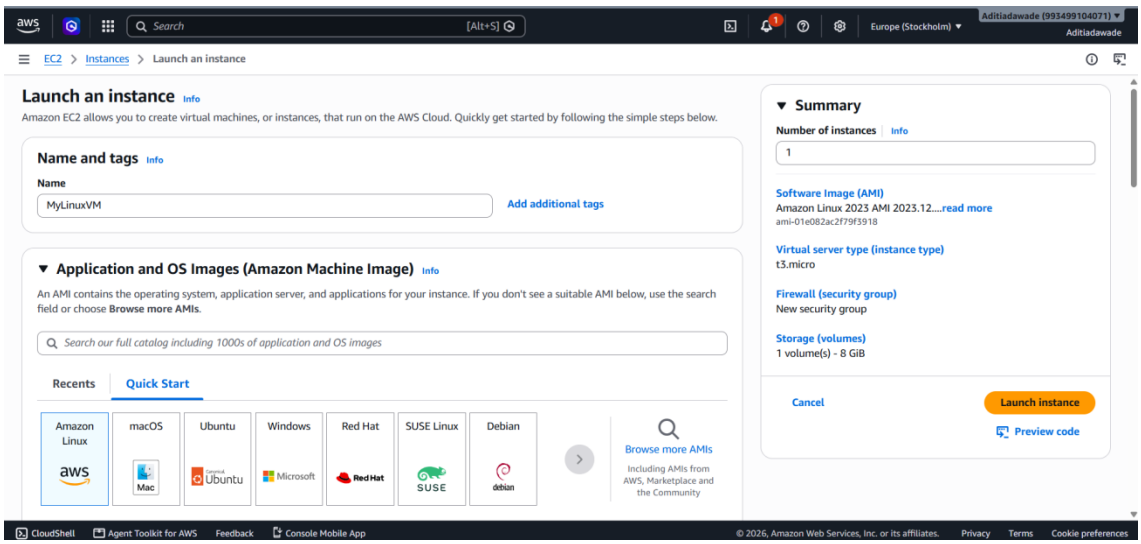
Step 1:log in to AWS

Step 2:Go to E2C Dashborad

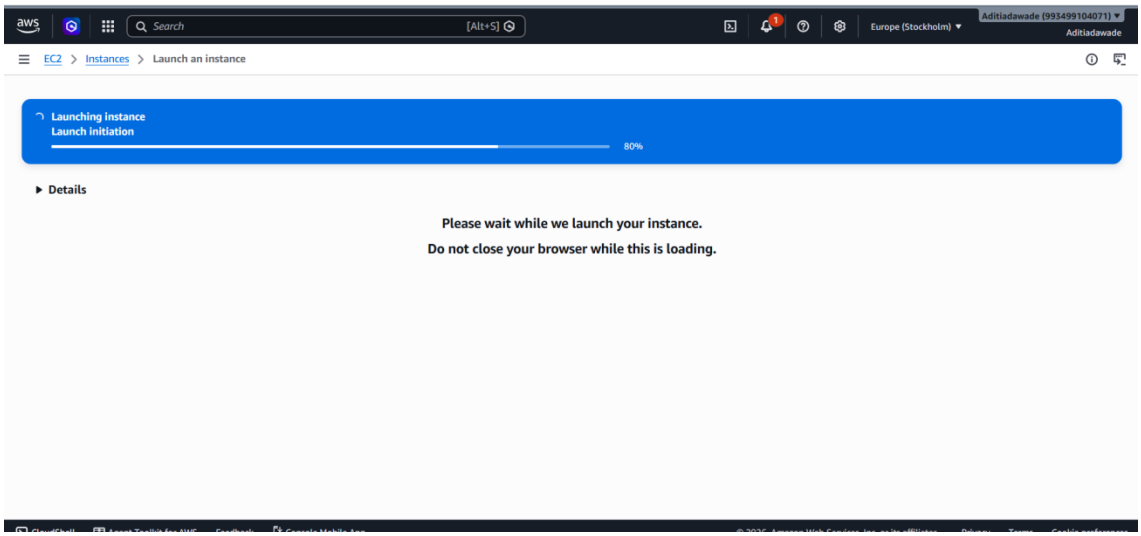


Step 3: Launch a new instance

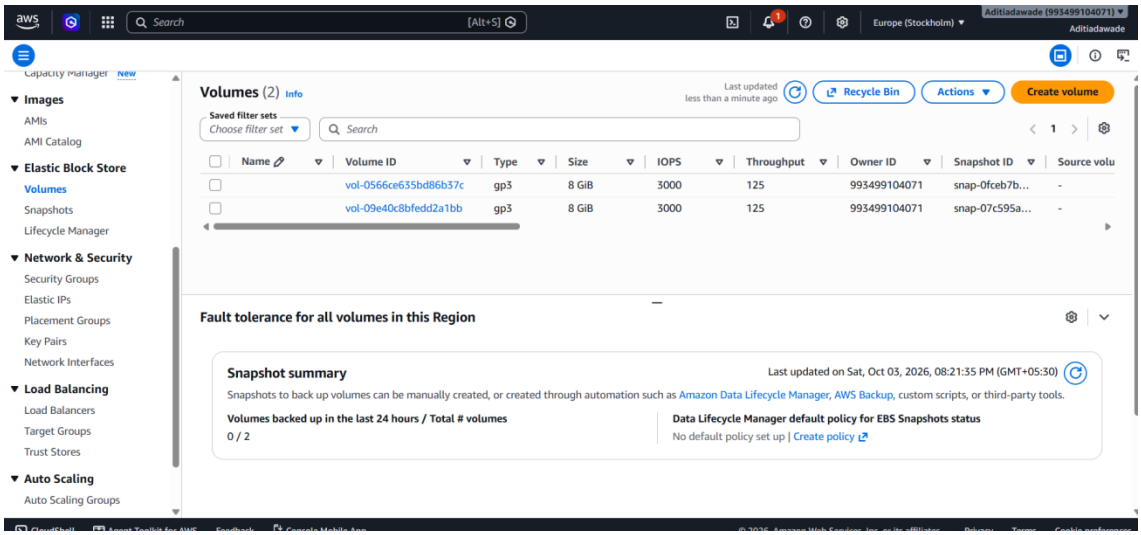
Step 4: configure instance details

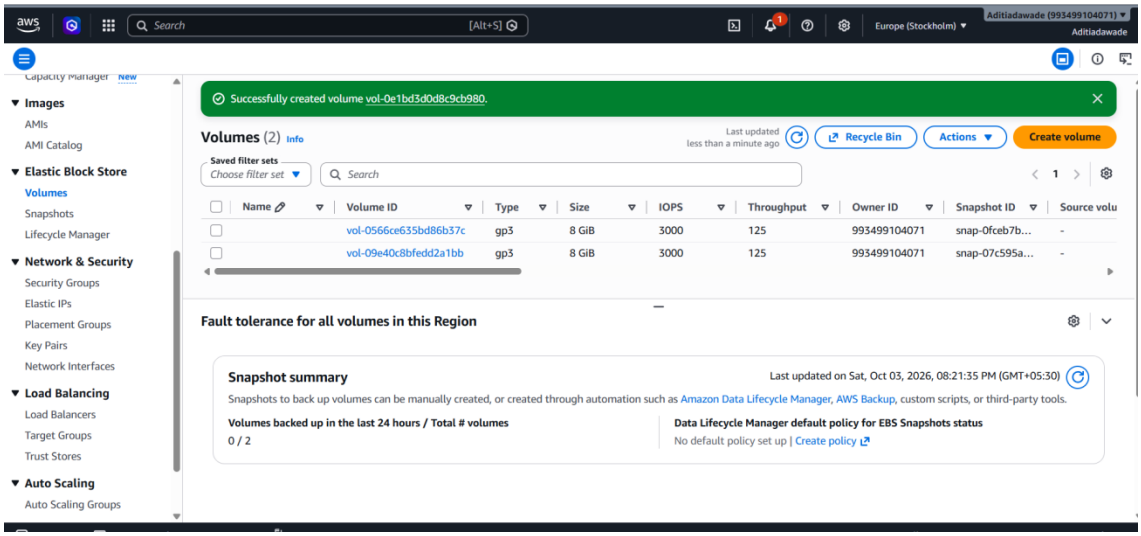


Step 6: Launch Instance

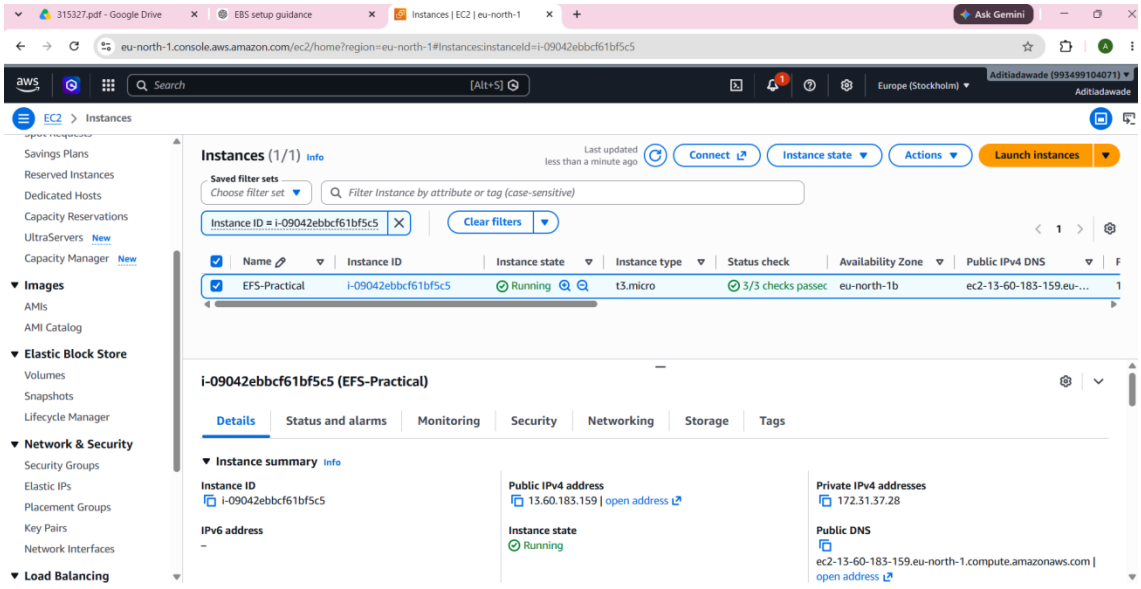


Step 7: Create an EBS Volume

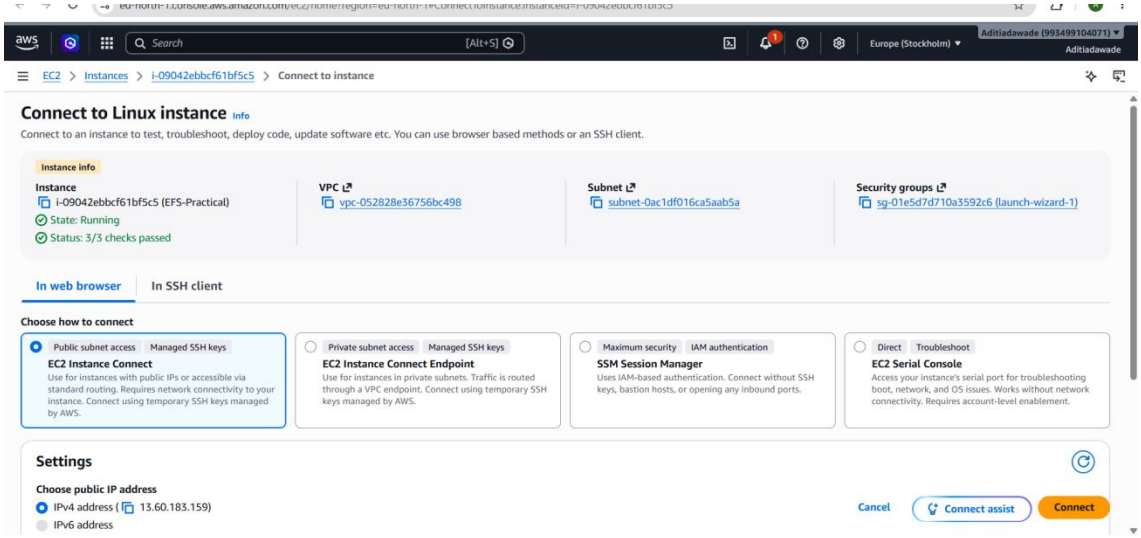




Step 8: Attach Volume to EC2 Instance



Step 9: Connect to your EC2 Instance



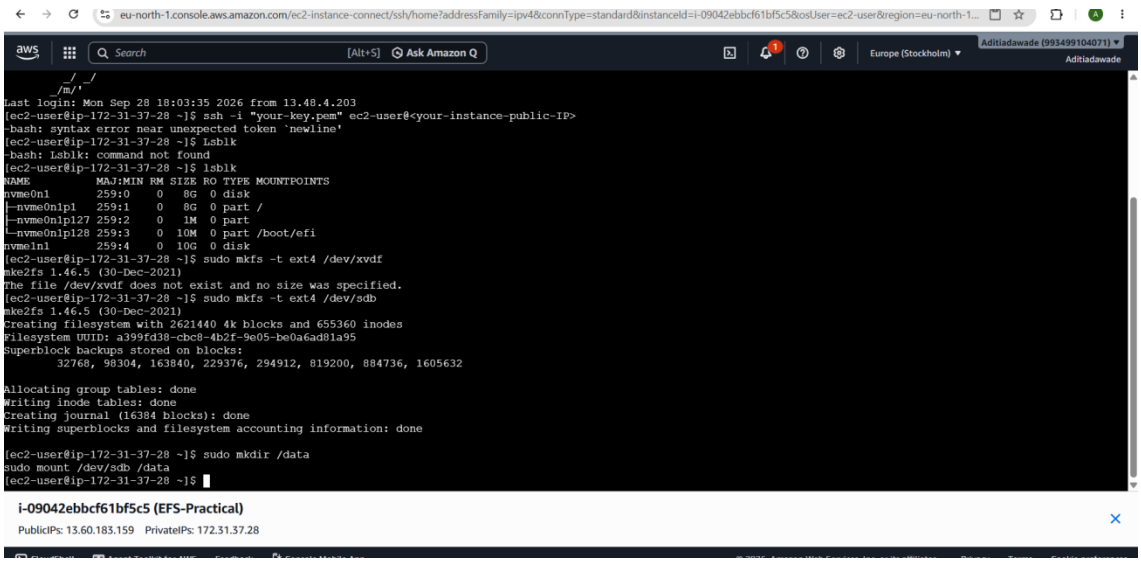
Step 10: List Attached Storage

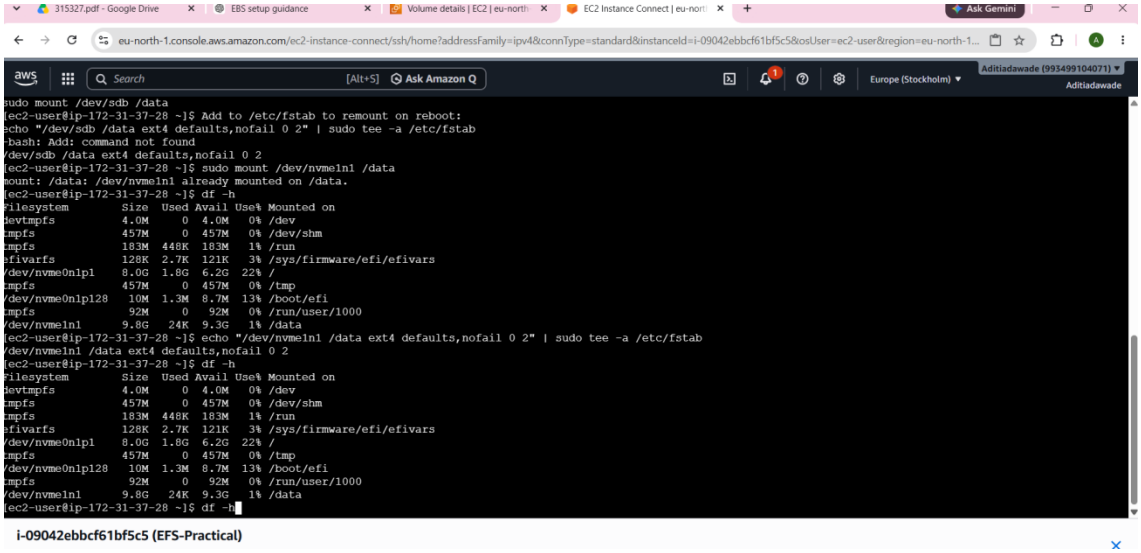
Step 11: Format the Volume

Step 12: Mount the Volume

Step 13: Make the Mount Persistent

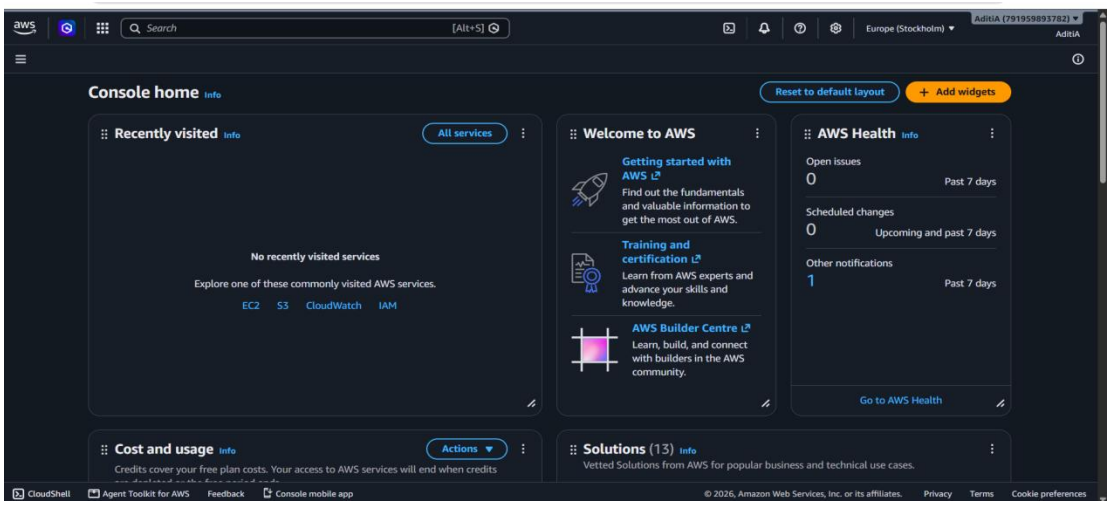
Step 14: Verify Everything

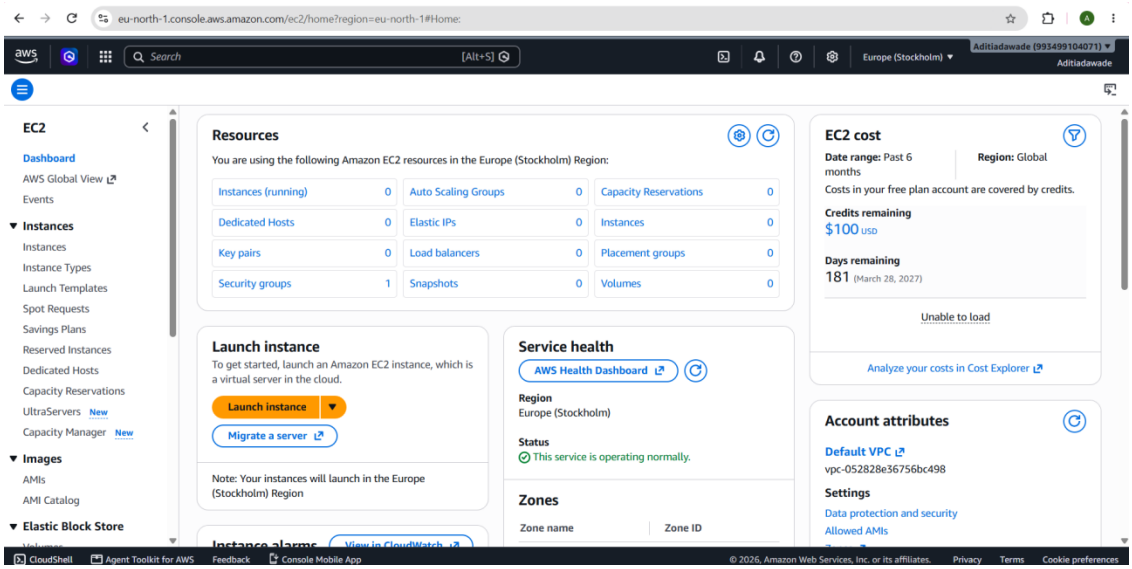


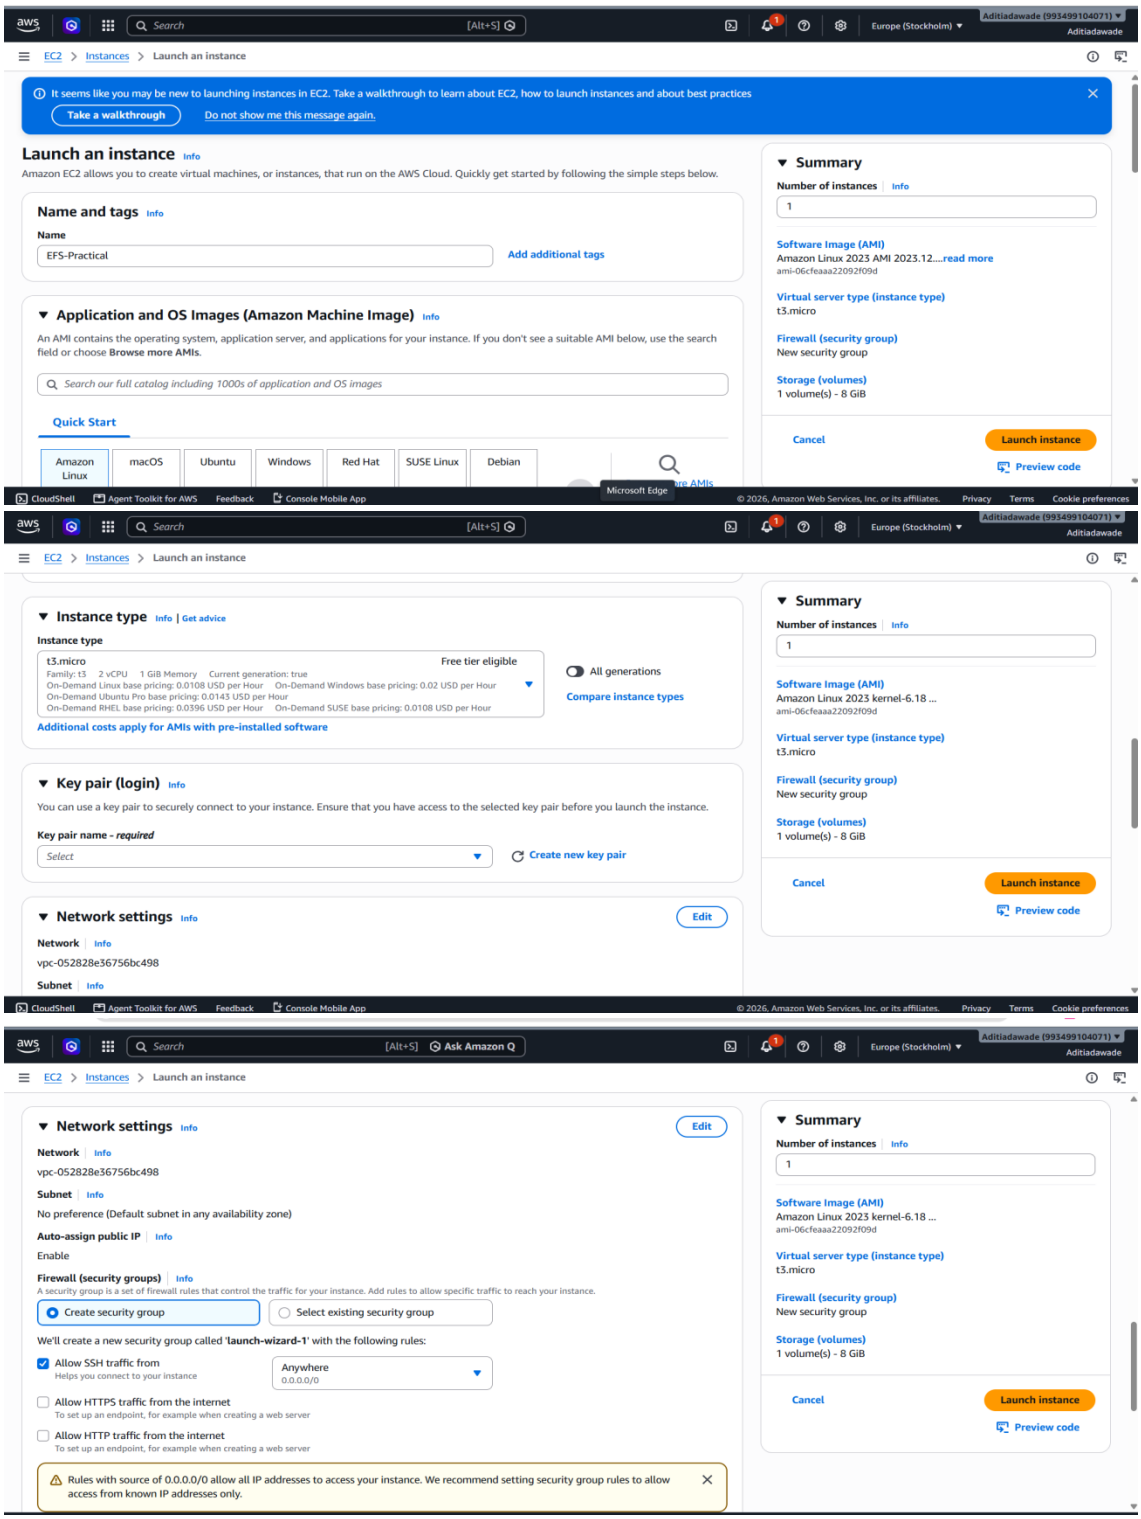

315327.pdf - Google Drive
EBS setup guidance
Volume details | EC2 | eu-north
EC2 Instance Connect | eu-north
Ask Gemini
eu-north-1.console.aws.amazon.com/ec2-instance-connect/ssh/home?addressFamily=ipv4&connType=standard&instanceId=i-09042ebbcf61bf5c5&osUser=ec2-user&region=eu-north-1...
aws
Search
[Alt+S]
Ask Amazon Q
Europe (Stockholm)
Aditiadawade (993499104071)
Aditiadawade
sudo mount /dev/sdb /data
[ec2-user@ip-172-31-37-28 ~]$ Add to /etc/fstab to remount on reboot:
echo "/dev/sdb /data ext4 defaults,nofail 0 2" | sudo tee -a /etc/fstab
-bash: Add: command not found
/dev/sdb /data ext4 defaults,nofail 0 2
[ec2-user@ip-172-31-37-28 ~]$ sudo mount /dev/nvme1n1 /data
mount: /data: /dev/nvme1n1 already mounted on /data.
[ec2-user@ip-172-31-37-28 ~]$ df -h
Filesystem Size Used Avail Use% Mounted on
devtmpfs 4.0M 0 4.0M 0% /dev
tmpfs 457M 0 457M 0% /dev/shm
tmpfs 183M 448K 183M 1% /run
efivarfs 128K 2.7K 121K 3% /sys/firmware/efi/efivars
/dev/nvme0n1p1 8.0G 1.8G 6.2G 22% /
tmpfs 457M 0 457M 0% /tmp
/dev/nvme0n1p128 10M 1.3M 8.7M 13% /boot/efi
tmpfs 92M 0 92M 0% /run/user/1000
/dev/nvme1n1 9.8G 24K 9.3G 1% /data
[ec2-user@ip-172-31-37-28 ~]$ echo "/dev/nvme1n1 /data ext4 defaults,nofail 0 2" | sudo tee -a /etc/fstab
/dev/nvme1n1 /data ext4 defaults,nofail 0 2
[ec2-user@ip-172-31-37-28 ~]$ df -h
Filesystem Size Used Avail Use% Mounted on
devtmpfs 4.0M 0 4.0M 0% /dev
tmpfs 457M 0 457M 0% /dev/shm
tmpfs 183M 448K 183M 1% /run
efivarfs 128K 2.7K 121K 3% /sys/firmware/efi/efivars
/dev/nvme0n1p1 8.0G 1.8G 6.2G 22% /
tmpfs 457M 0 457M 0% /tmp
/dev/nvme0n1p128 10M 1.3M 8.7M 13% /boot/efi
tmpfs 92M 0 92M 0% /run/user/1000
/dev/nvme1n1 9.8G 24K 9.3G 1% /data
[ec2-user@ip-172-31-37-28 ~]$ df -h
i-09042ebbcf61bf5c5 (EFS-Practical)

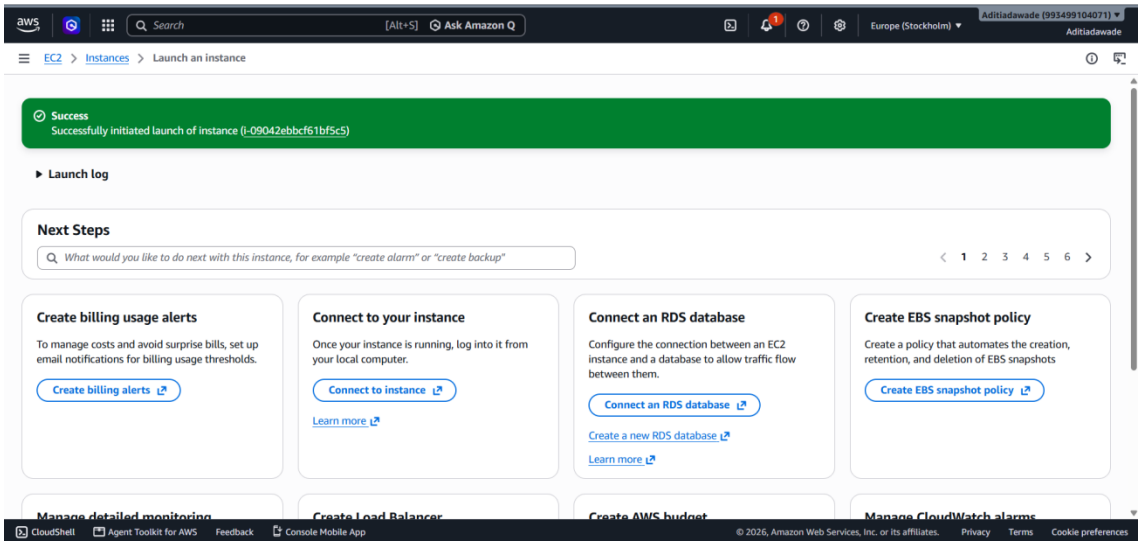
Practical no. 7

Step1: Launch an ec2 instance

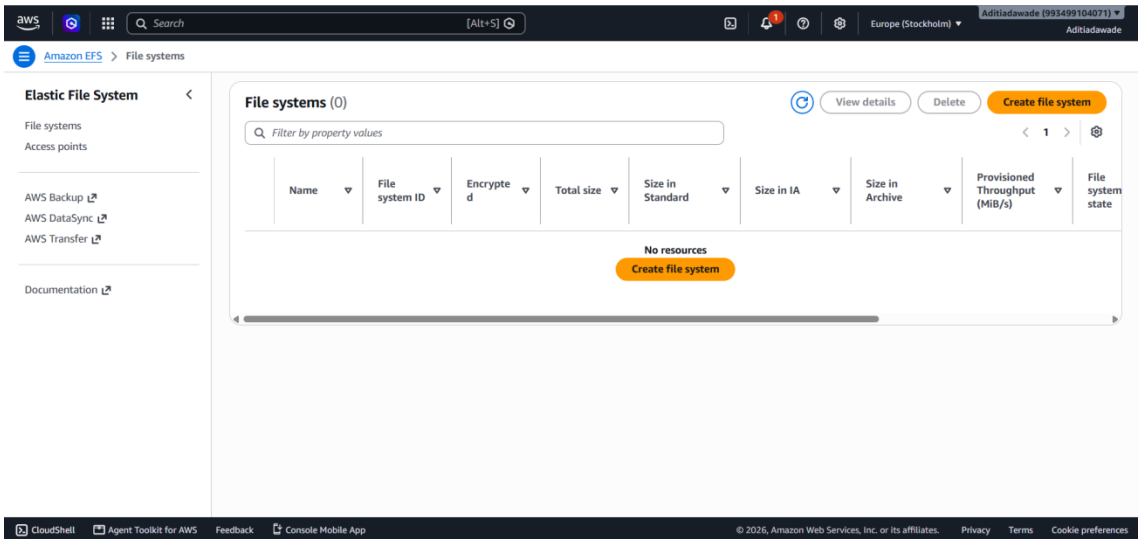


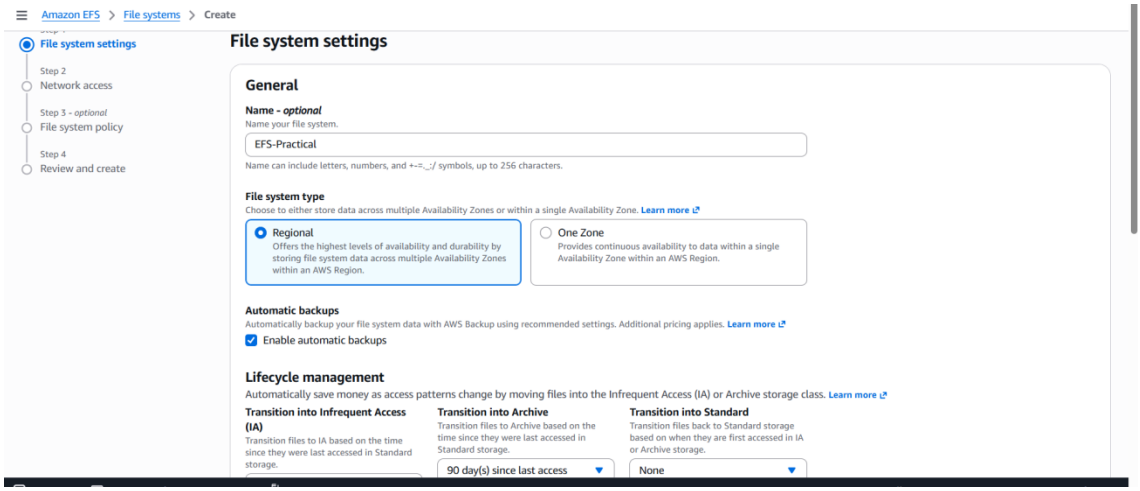


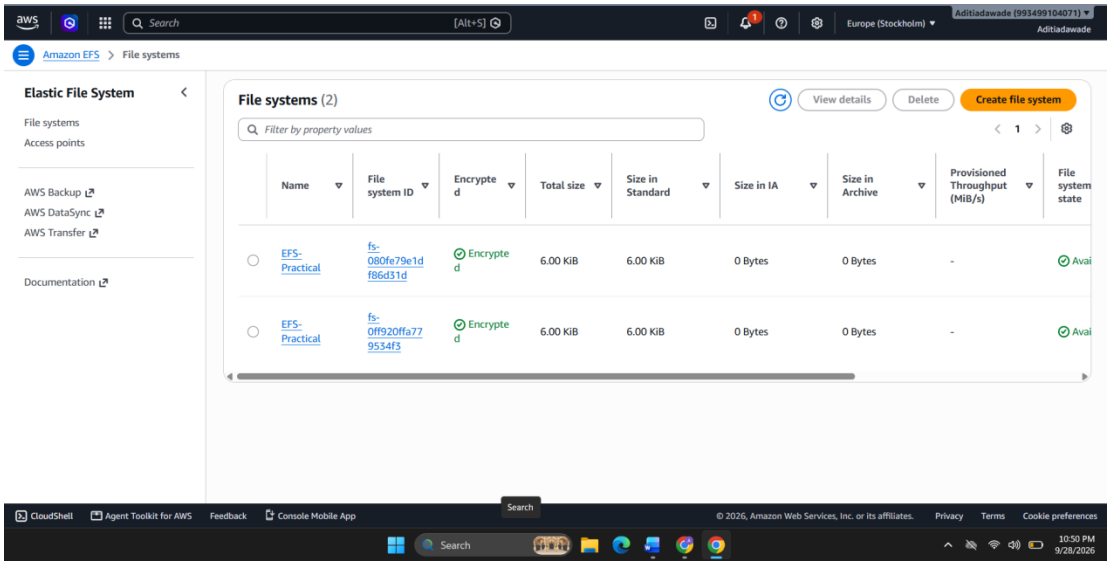

Launch an instance
Summary
Name and tags
EFS-Practical
Application and OS Images (Amazon Machine Image)
Quick Start
Amazon Linux
macOS
Ubuntu
Windows
Red Hat
SUSE Linux
Debian
Instance type
t3.micro
Free tier eligible
Key pair (login)
Key pair name - required
Select
Create new key pair
Network settings
Edit
Network
vpc-052828e36756bc498
Subnet
No preference (Default subnet in any availability zone)
Auto-assign public IP
Enable
Firewall (security groups)
Create security group
Select existing security group
We'll create a new security group called 'launch-wizard-1' with the following rules:
Allow SSH traffic from
Anywhere
0.0.0.0/0
Allow HTTPS traffic from the internet
Allow HTTP traffic from the internet
Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.
Number of instances
Software Image (AMI)
Virtual server type (instance type)
t3.micro
Firewall (security group)
New security group
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Launch instance
Preview code

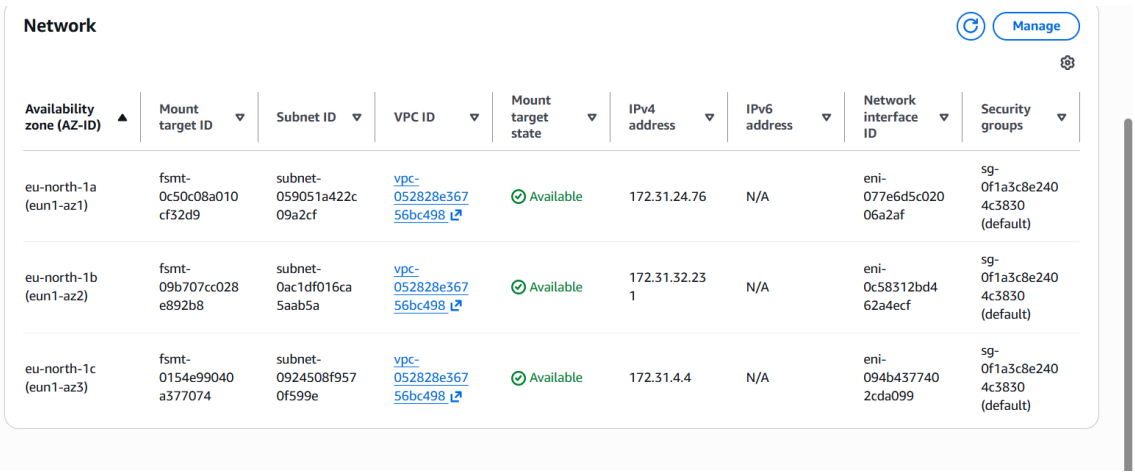


Step2: Create an EFS File system

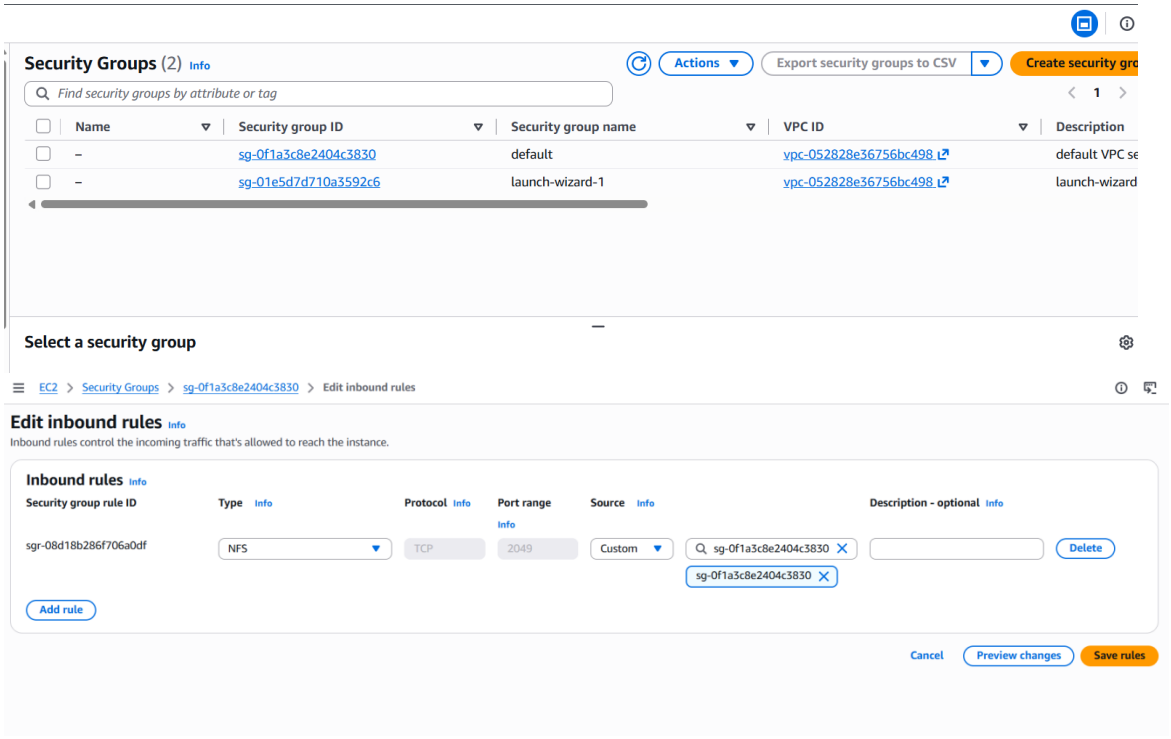


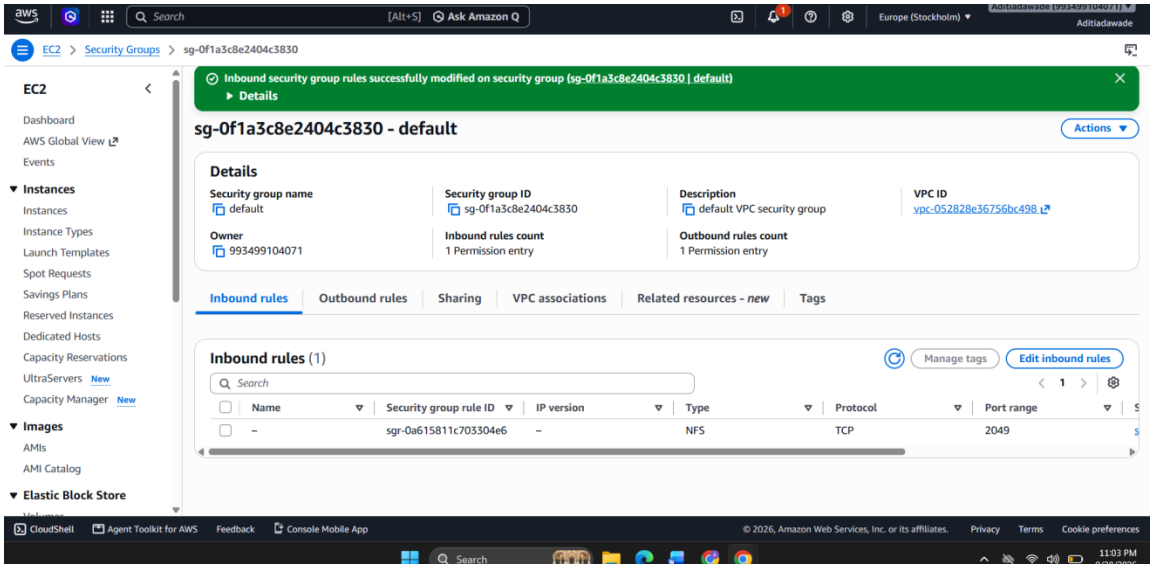

Amazon EFS > File systems > Create
File system settings
Step 2
Network access
Step 3 - optional
File system policy
Step 4
Review and create
File system settings
General
Name - optional
Name your file system.
EFS-Practical
Name can include letters, numbers, and +-=._:/ symbols, up to 256 characters.
File system type
Choose to either store data across multiple Availability Zones or within a single Availability Zone. Learn more
Regional
Offers the highest levels of availability and durability by storing file system data across multiple Availability Zones within an AWS Region.
One Zone
Provides continuous availability to data within a single Availability Zone within an AWS Region.
Automatic backups
Automatically backup your file system data with AWS Backup using recommended settings. Additional pricing applies. Learn more
Enable automatic backups
Lifecycle management
Automatically save money as access patterns change by moving files into the Infrequent Access (IA) or Archive storage class. Learn more
Transition into Infrequent Access (IA)
Transition files to IA based on the time since they were last accessed in Standard storage.
Transition into Archive
Transition files to Archive based on the time since they were last accessed in Standard storage.
90 day(s) since last access
Transition into Standard
Transition files back to Standard storage based on when they are first accessed in IA or Archive storage.
None

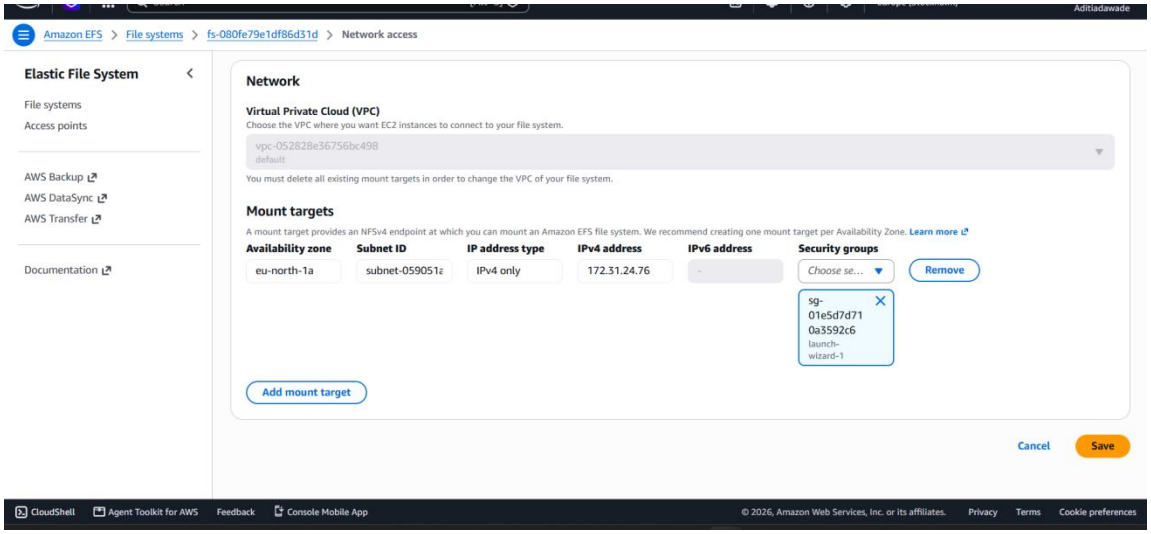

Search
[Alt+S]
Europe (Stockholm)
Aditiadawade (993499104071)
Aditiadawade
Amazon EFS > File systems
Elastic File System
File systems
Access points
AWS Backup
AWS DataSync
AWS Transfer
Documentation
File systems (2)
View details
Delete
Create file system
Filter by property values
Name
File system ID
Encrypted
Total size
Size in Standard
Size in IA
Size in Archive
Provisioned Throughput (MiB/s)
File system state
EFS-Practical
fs-080fe79e1df86d31d
Encrypted
6.00 KiB
6.00 KiB
0 Bytes
0 Bytes
EFS-Practical
fs-0ff920ffa779534f3
Encrypted
6.00 KiB
6.00 KiB
0 Bytes
0 Bytes
CloudShell
Agent Toolkit for AWS
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
10:50 PM
9/28/2026

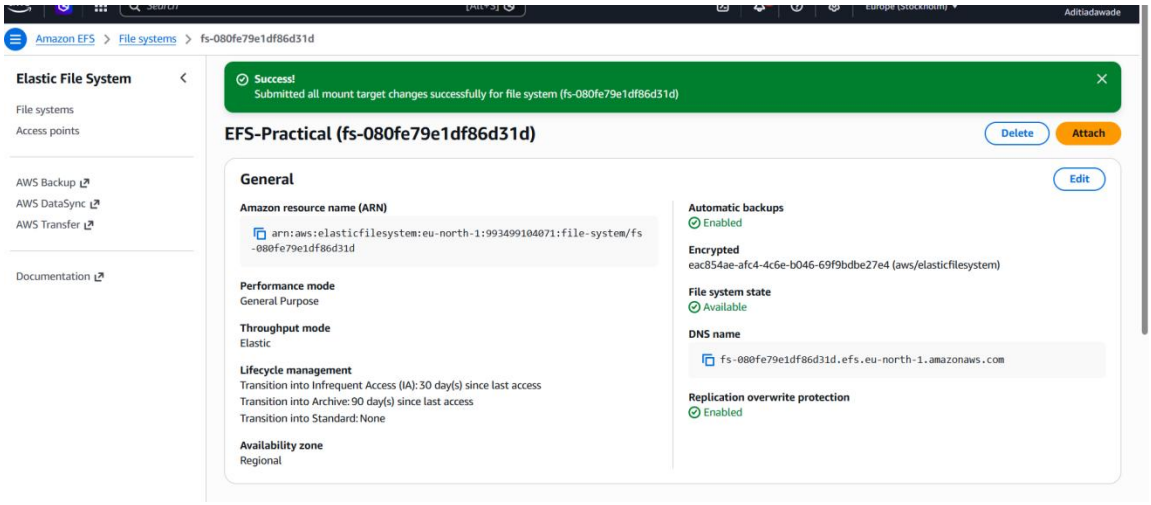

Network
Manage
Availability zone (AZ-ID)
Mount target ID
Subnet ID
VPC ID
Mount target state
IPv4 address
IPv6 address
Network interface ID
Security groups
eu-north-1a (eun1-az1)
fsmt-0c50c08a010cf32d9
subnet-059051a422c09a2cf
vpc-052828e36756bc498
Available
172.31.24.76
N/A
eni-077e6d5c02006a2af
sg-0f1a3c8e2404c3830 (default)
eu-north-1b (eun1-az2)
fsmt-09b707cc028e892b8
subnet-0ac1df016ca5aab5a
vpc-052828e36756bc498
Available
172.31.32.231
N/A
eni-0c58312bd462a4ecf
sg-0f1a3c8e2404c3830 (default)
eu-north-1c (eun1-az3)
fsmt-0154e990400a377074
subnet-0924508f9570f599e
vpc-052828e36756bc498
Available
172.31.4.4
N/A
eni-094b4377402cda099
sg-0f1a3c8e2404c3830 (default)

Step 3:Configure Security Groups





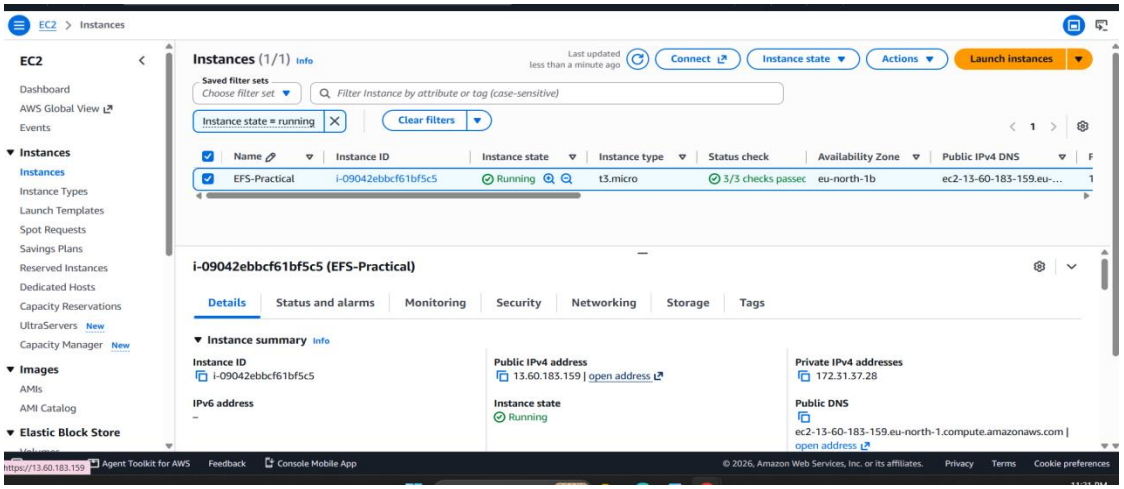


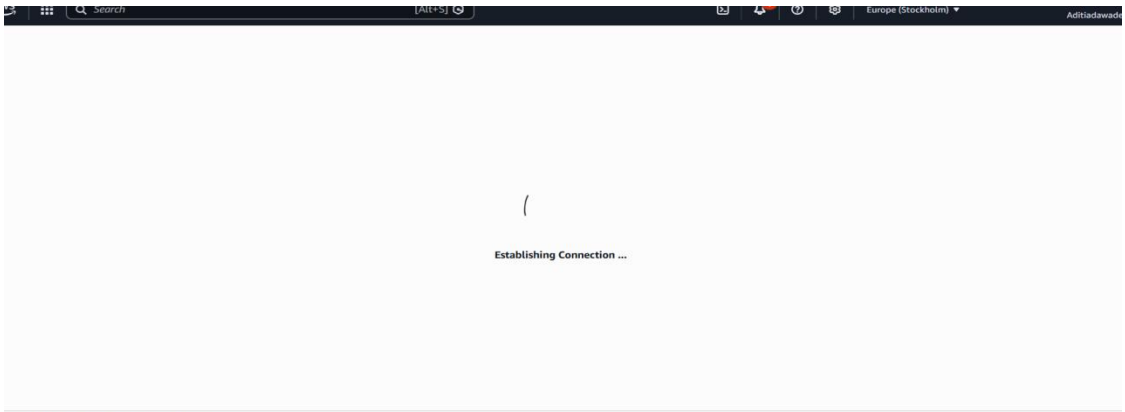


Step4: Connect to EC2 instance vis SSH

Step5: Install NFS utilites on EC2

Step6: Create a mount Directory on EC2





09042ebbcf61bf5c5 (EFS-Practical)

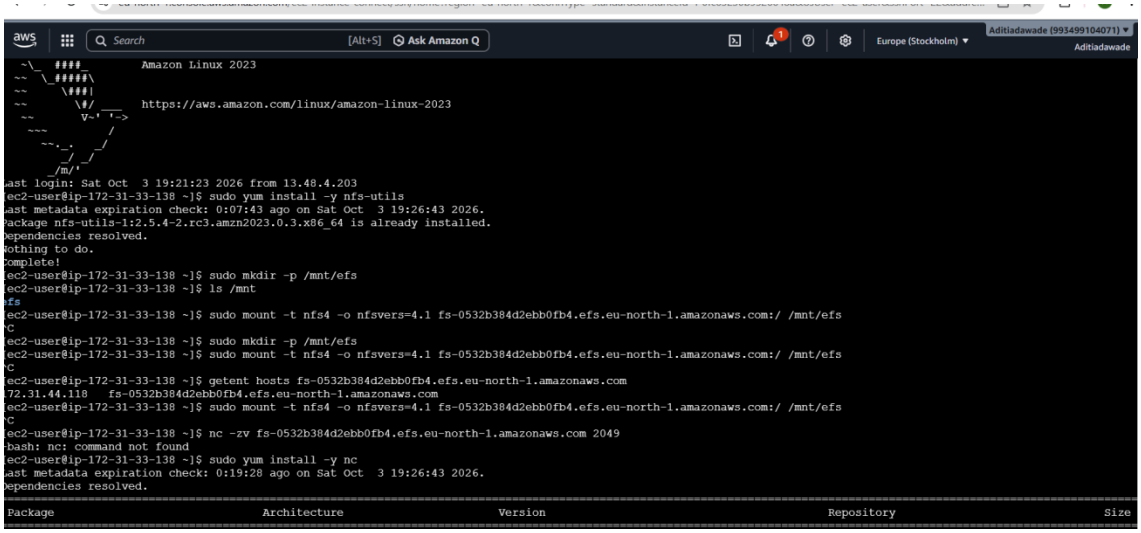
```
  ~~  \_#####\
  ~~     \###|
  ~~       \#/ ___    https://aws.amazon.com/linux/amazon-linux-2023
   ~~       V~' '->
    ~~~         /
      ~~._.   _/
         _/ _/
       _/m/'
[ec2-user@ip-172-31-33-138 ~]$ ssh -i "your-key.pem" ec2-user@<Public-IP>
-bash: syntax error near unexpected token `newline'
[ec2-user@ip-172-31-33-138 ~]$ ssh -i "C:\path\to\your-key.pem" ec2-user@ssh51.20.3.193
Warning: Identity file C:\path\to\your-key.pem not accessible: No such file or directory.
ssh: Could not resolve hostname ssh51.20.3.193: Name or service not known
[ec2-user@ip-172-31-33-138 ~]$ sudo yum install -y nfs-utils

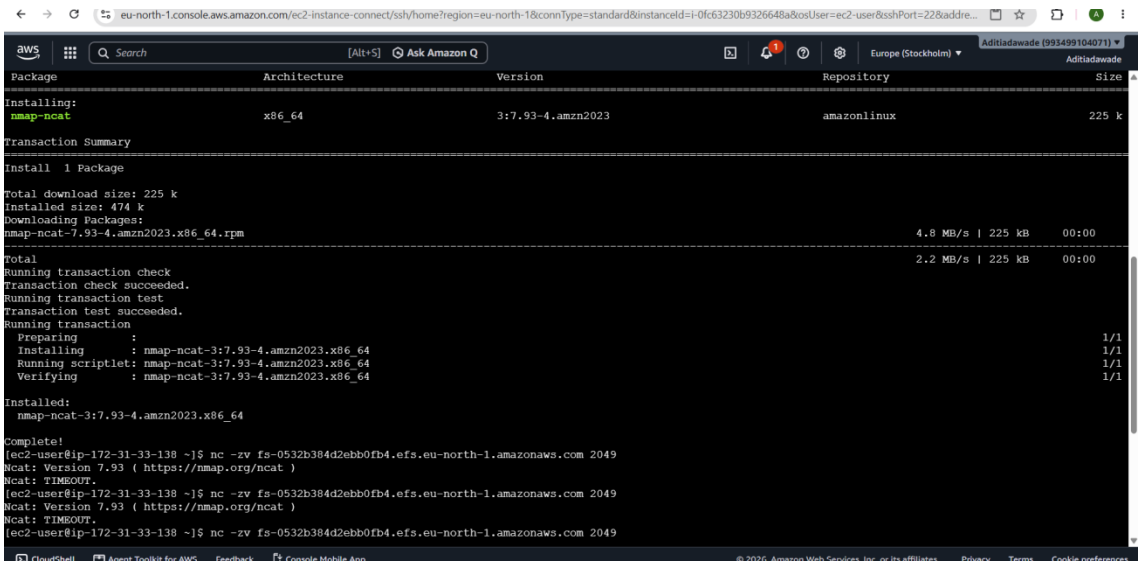
Amazon Linux 2023 Kernel Livepatch repository
Package nfs-utils-1:2.5.4-2.rc3.amzn2023.0.3.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-33-138 ~]$ sudo mkdir -p /mnt/efs
[ec2-user@ip-172-31-33-138 ~]$ ls /mnt
efs
[ec2-user@ip-172-31-33-138 ~]$
```

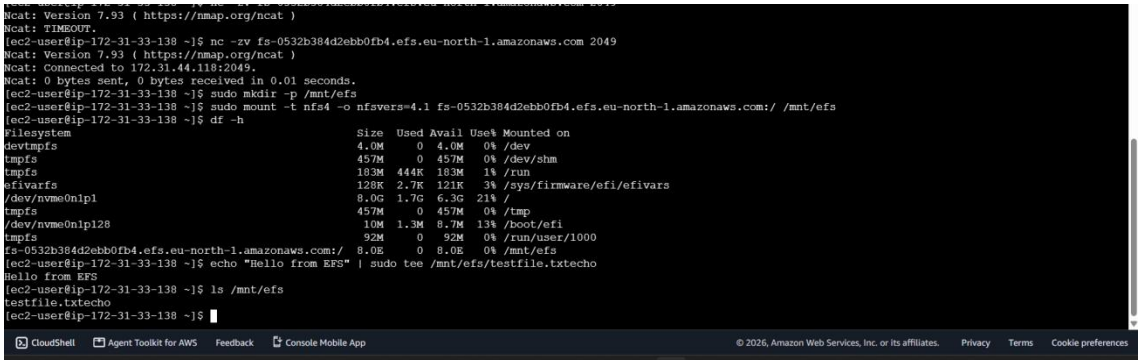
Step 7: Mount the EFS File System

Step 8: Verify the mount

Step 9: Test the storage by creating a file

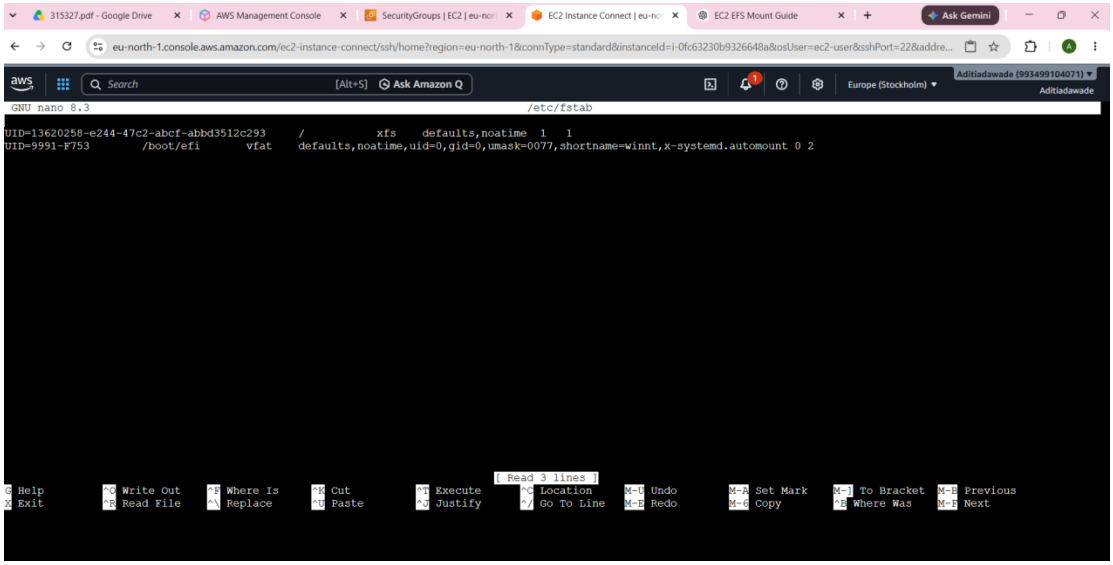


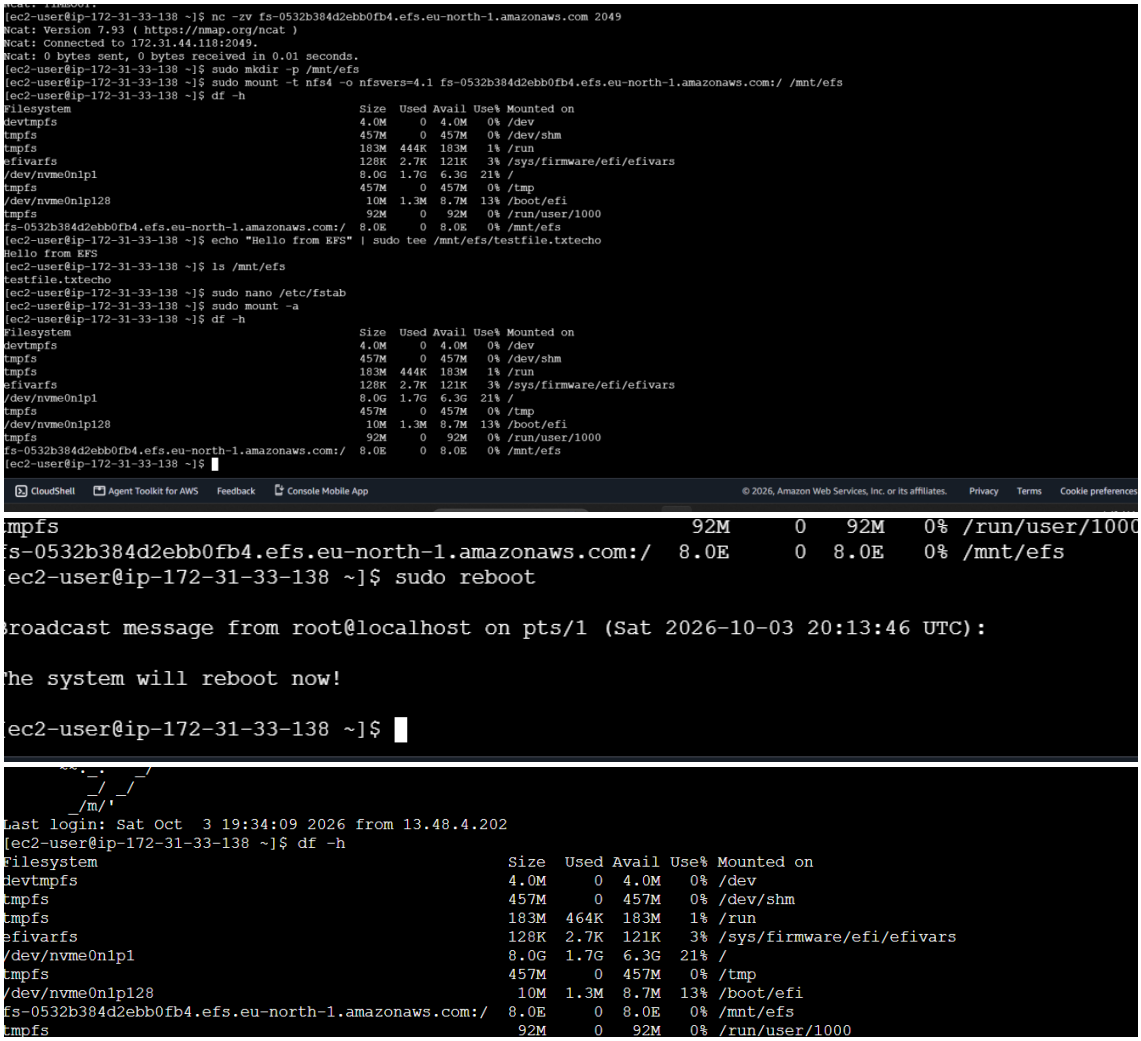




Step 10: Make the mount persistent

Step 11: Reboot and confirm persistent



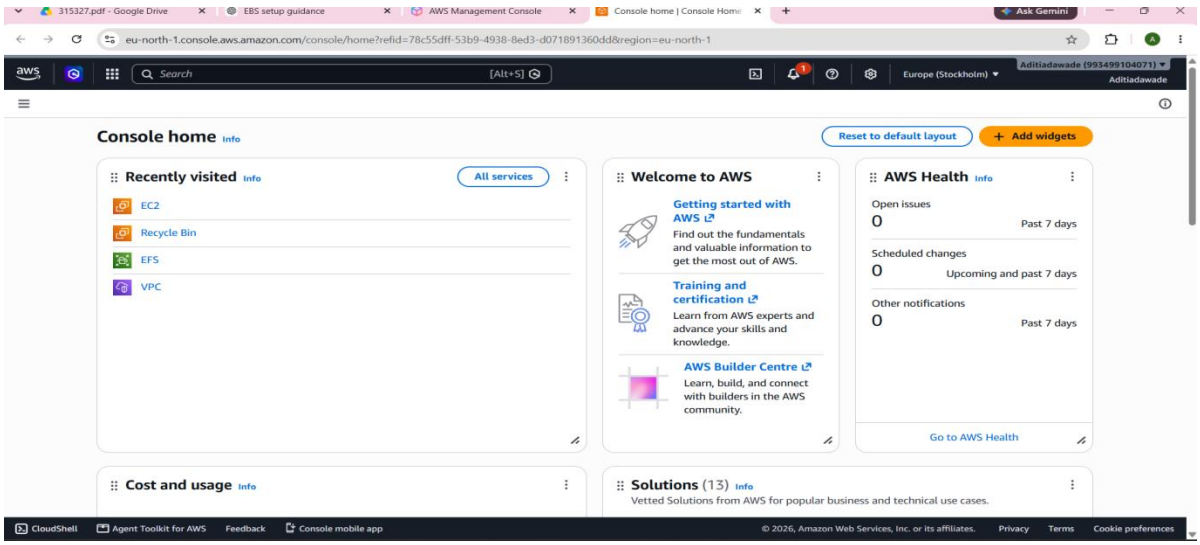

[ec2-user@ip-172-31-33-138 ~]$ nc -zv fs-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com 2049
Ncat: Version 7.93 (https://nmap.org/ncat)
Ncat: Connected to 172.31.44.118:2049.
Ncat: 0 bytes sent, 0 bytes received in 0.01 seconds.
[ec2-user@ip-172-31-33-138 ~]$ sudo mkdir -p /mnt/efs
[ec2-user@ip-172-31-33-138 ~]$ sudo mount -t nfs4 -o nfsvers=4.1 fs-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com:/ /mnt/efs
[ec2-user@ip-172-31-33-138 ~]$ df -h
Filesystem Size Used Avail Use% Mounted on
devtmpfs 4.0M 0 4.0M 0% /dev
tmpfs 457M 0 457M 0% /dev/shm
tmpfs 183M 444K 183M 1% /run
efivarfs 128K 2.7K 121K 3% /sys/firmware/efi/efivars
/dev/nvme0n1p1 8.0G 1.7G 6.3G 21% /
tmpfs 457M 0 457M 0% /tmp
/dev/nvme0n1p128 10M 1.3M 8.7M 13% /boot/efi
tmpfs 92M 0 92M 0% /run/user/1000
fs-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com:/ 8.0E 0 8.0E 0% /mnt/efs
[ec2-user@ip-172-31-33-138 ~]$ echo "Hello from EFS" | sudo tee /mnt/efs/testfile.txtecho
Hello from EFS
[ec2-user@ip-172-31-33-138 ~]$ ls /mnt/efs
testfile.txtecho
[ec2-user@ip-172-31-33-138 ~]$ sudo nano /etc/fstab
[ec2-user@ip-172-31-33-138 ~]$ sudo mount -a
[ec2-user@ip-172-31-33-138 ~]$ df -h
Filesystem Size Used Avail Use% Mounted on
devtmpfs 4.0M 0 4.0M 0% /dev
tmpfs 457M 0 457M 0% /dev/shm
tmpfs 183M 444K 183M 1% /run
efivarfs 128K 2.7K 121K 3% /sys/firmware/efi/efivars
/dev/nvme0n1p1 8.0G 1.7G 6.3G 21% /
tmpfs 457M 0 457M 0% /tmp
/dev/nvme0n1p128 10M 1.3M 8.7M 13% /boot/efi
tmpfs 92M 0 92M 0% /run/user/1000
fs-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com:/ 8.0E 0 8.0E 0% /mnt/efs
[ec2-user@ip-172-31-33-138 ~]$
CloudShell
Agent Toolkit for AWS
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

mpfs 92M 0 92M 0% /run/user/1000
s-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com:/ 8.0E 0 8.0E 0% /mnt/efs
ec2-user@ip-172-31-33-138 ~]$ sudo reboot
roadcast message from root@localhost on pts/1 (Sat 2026-10-03 20:13:46 UTC):
he system will reboot now!
ec2-user@ip-172-31-33-138 ~]$

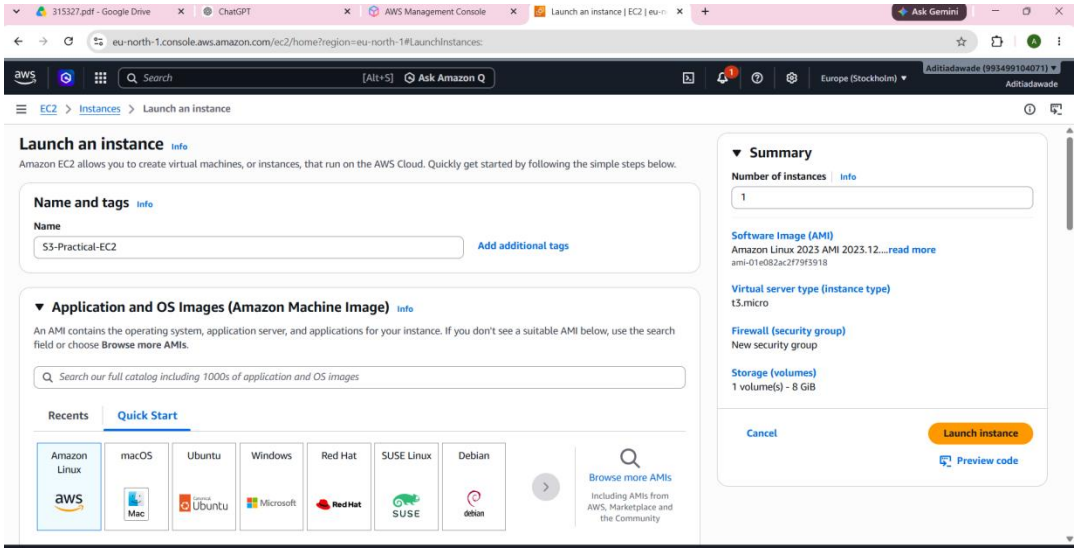
Last login: Sat Oct 3 19:34:09 2026 from 13.48.4.202
[ec2-user@ip-172-31-33-138 ~]$ df -h
Filesystem Size Used Avail Use% Mounted on
devtmpfs 4.0M 0 4.0M 0% /dev
tmpfs 457M 0 457M 0% /dev/shm
tmpfs 183M 464K 183M 1% /run
efivarfs 128K 2.7K 121K 3% /sys/firmware/efi/efivars
/dev/nvme0n1p1 8.0G 1.7G 6.3G 21% /
tmpfs 457M 0 457M 0% /tmp
/dev/nvme0n1p128 10M 1.3M 8.7M 13% /boot/efi
fs-0532b384d2ebb0fb4.efs.eu-north-1.amazonaws.com:/ 8.0E 0 8.0E 0% /mnt/efs
tmpfs 92M 0 92M 0% /run/user/1000

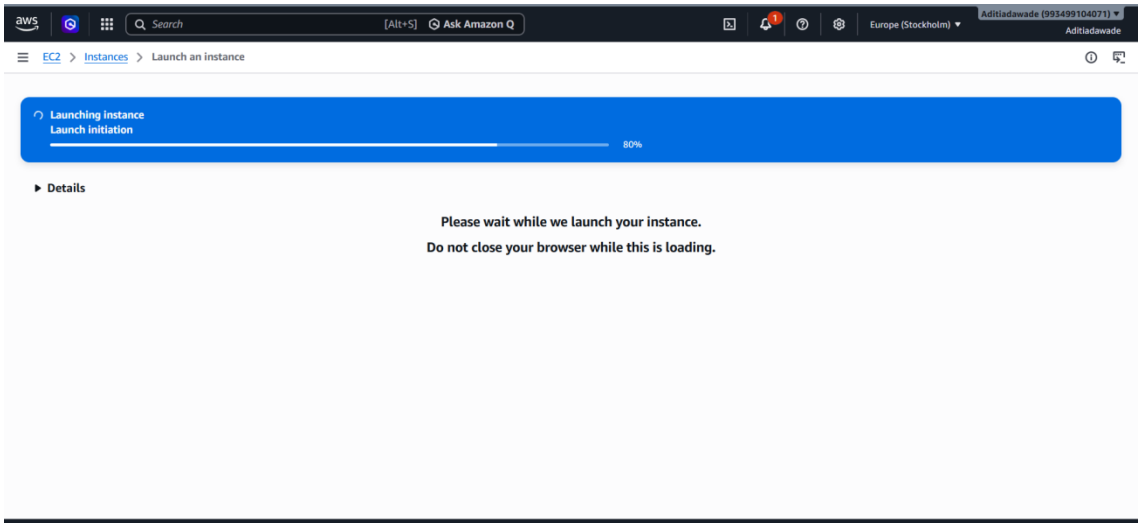
Practical no. 8

Step 1: Log in to AWS Console

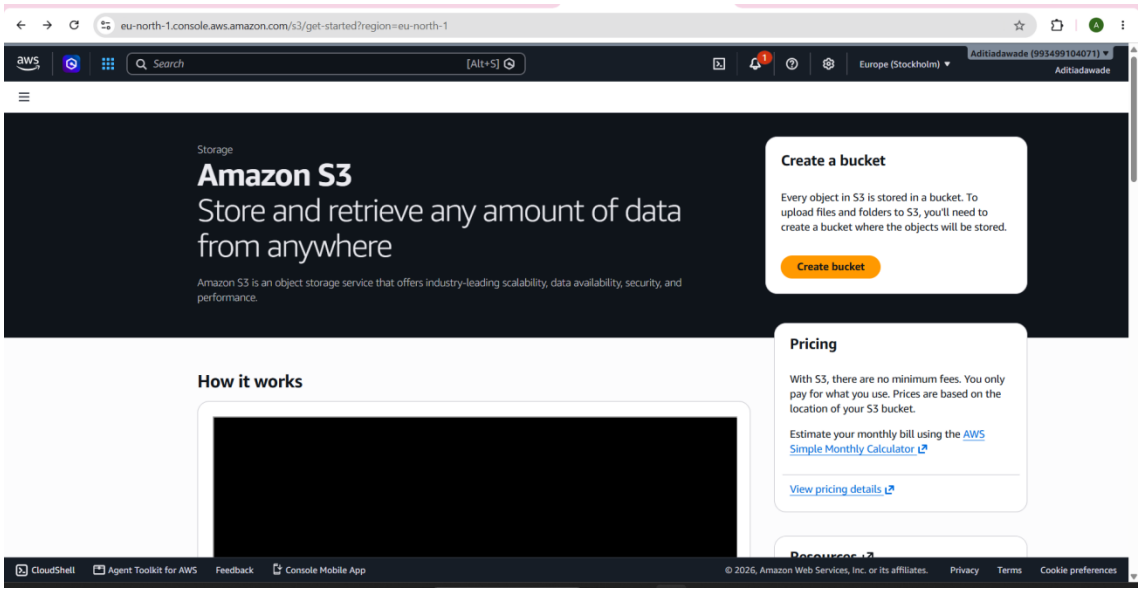


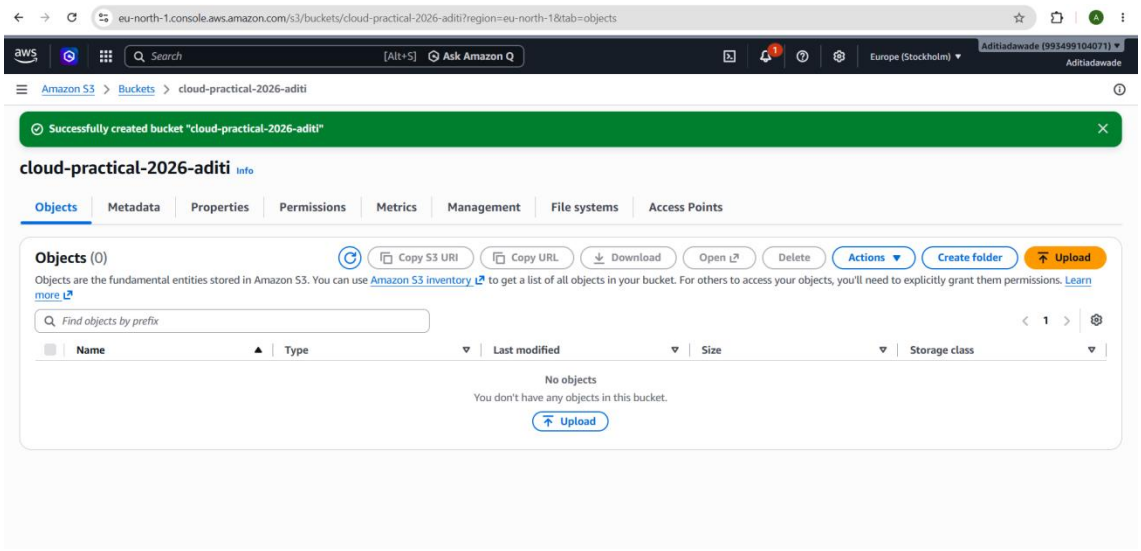
Step 2: Launch an EC2 Instance



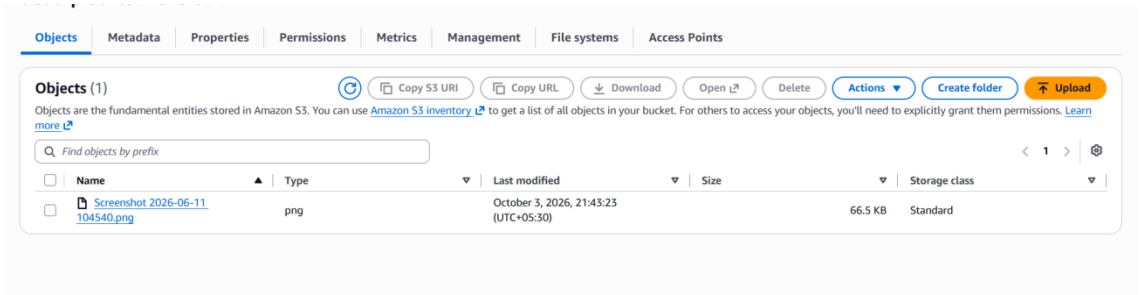


Step 3: Create s3 bucket

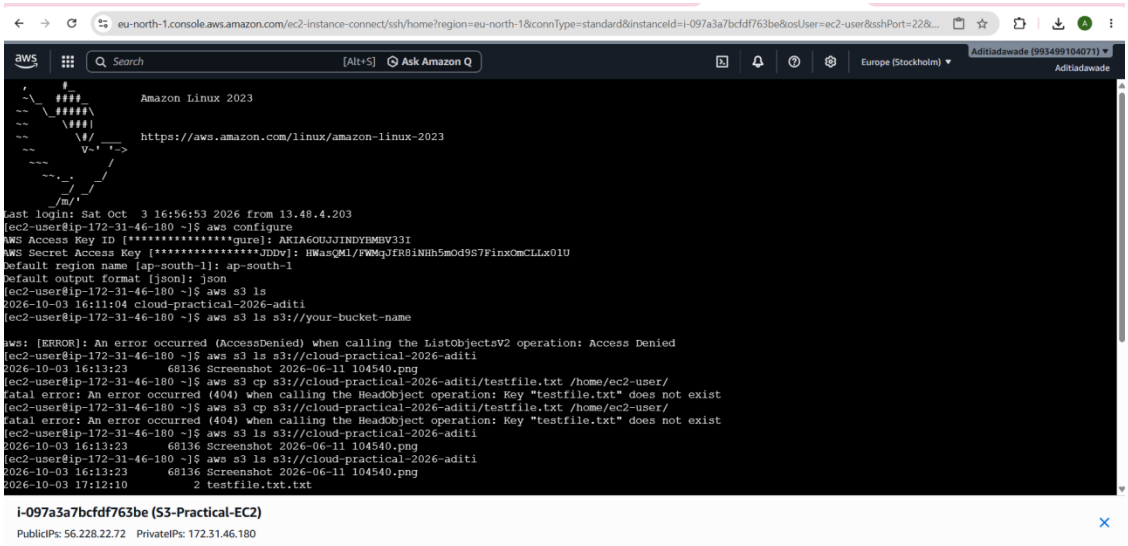




Step 4: Upload object to the S3 bucket



Step 7: Access S3 bucket from EC2

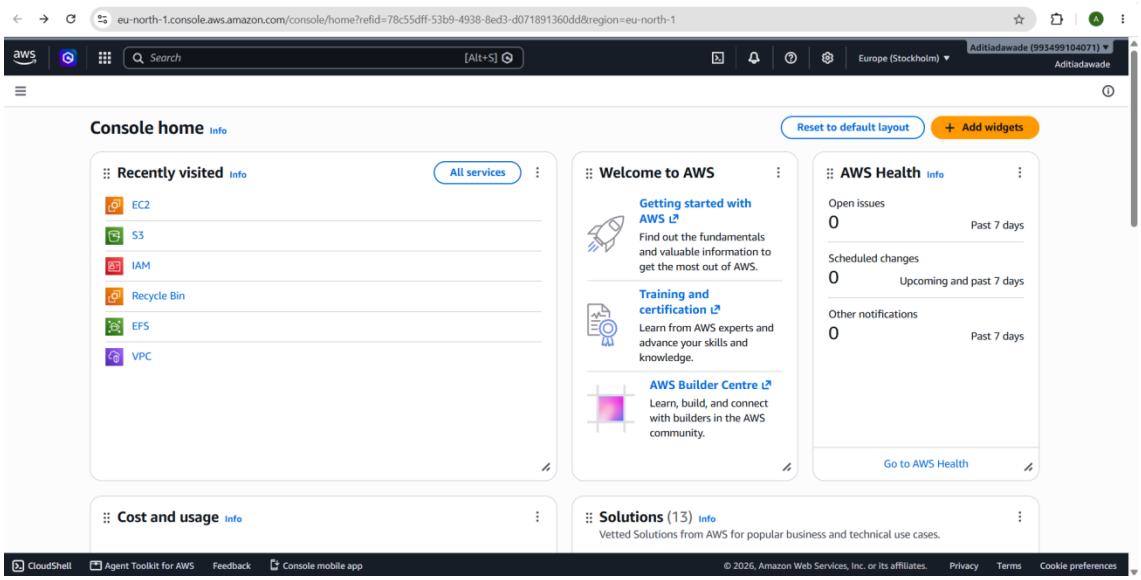


Step 8: Verify and clean up

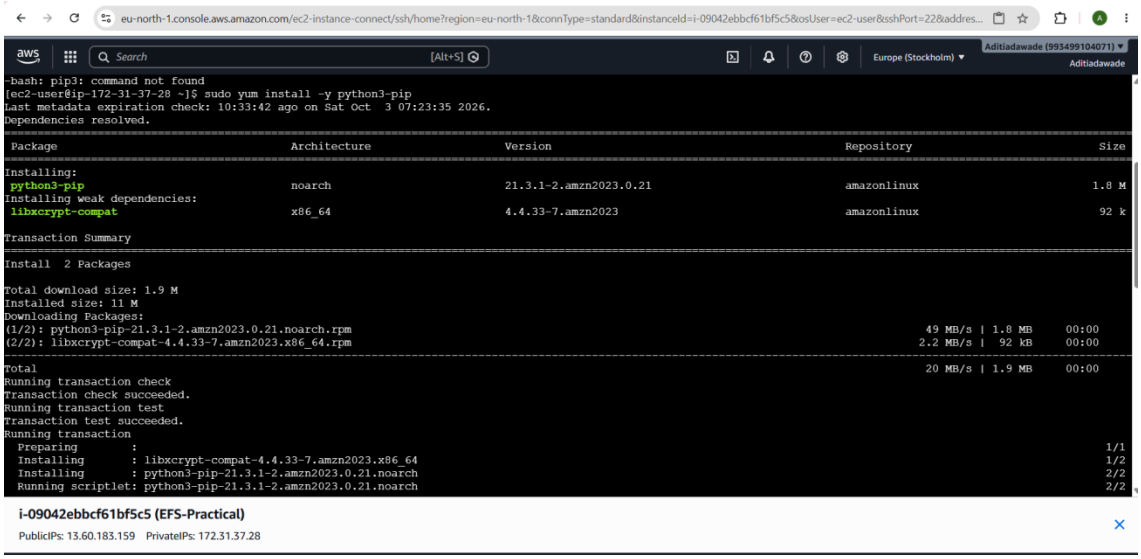


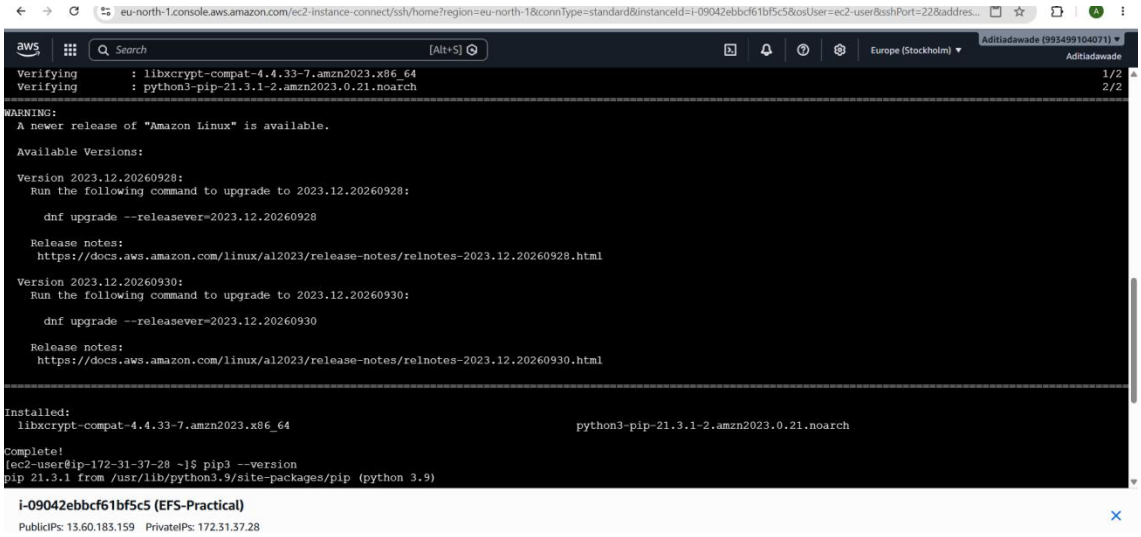
Practical no. 9

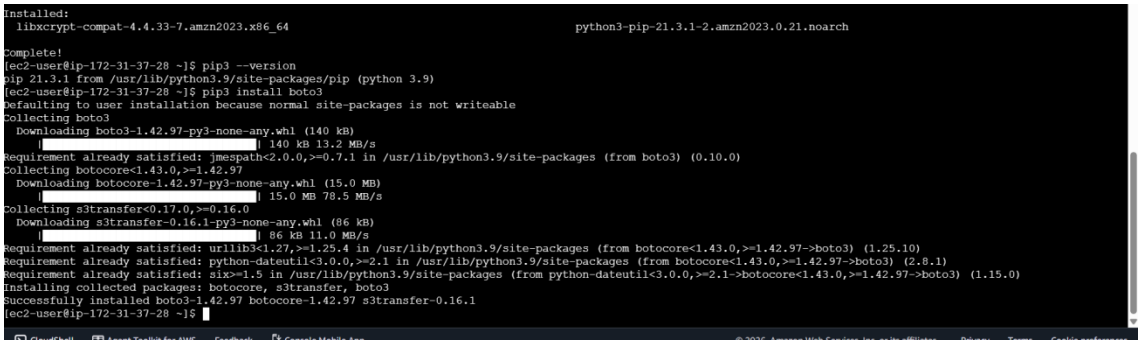
Step1: Set up AWS account and IAM



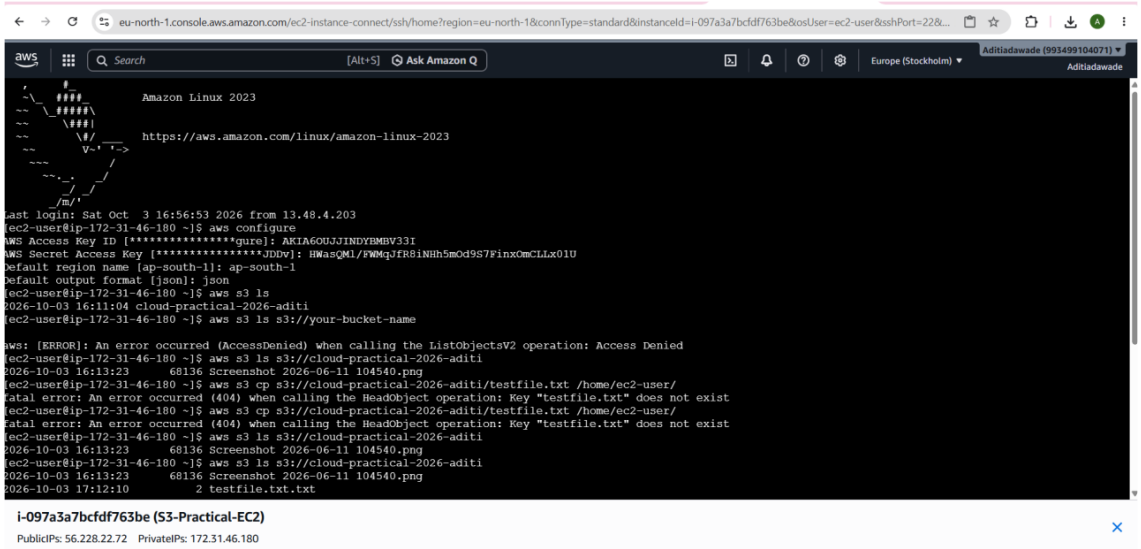
Step2: Install Required tools





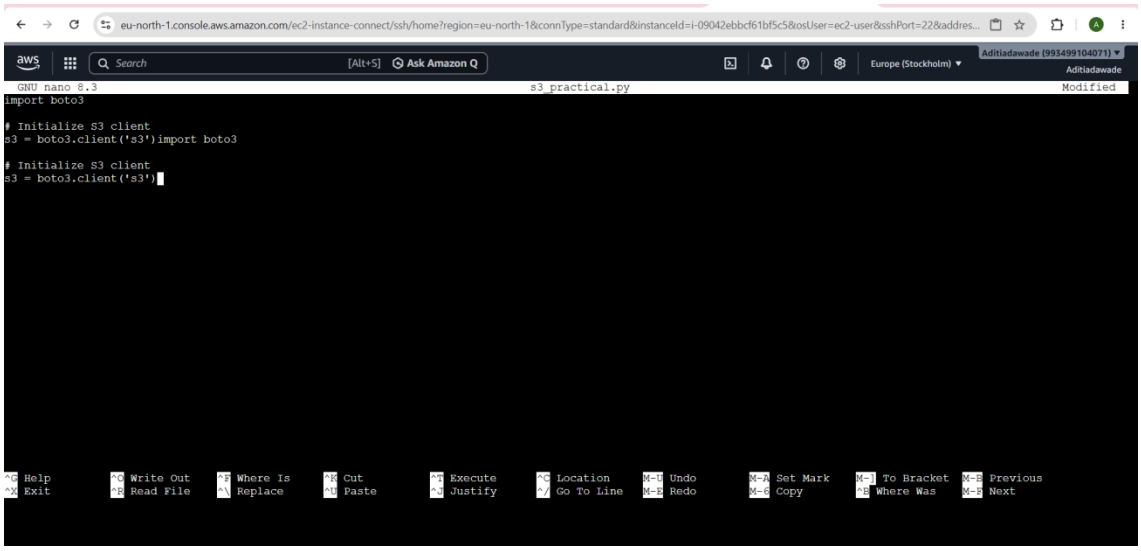


Step 3: Configure AWS credentials

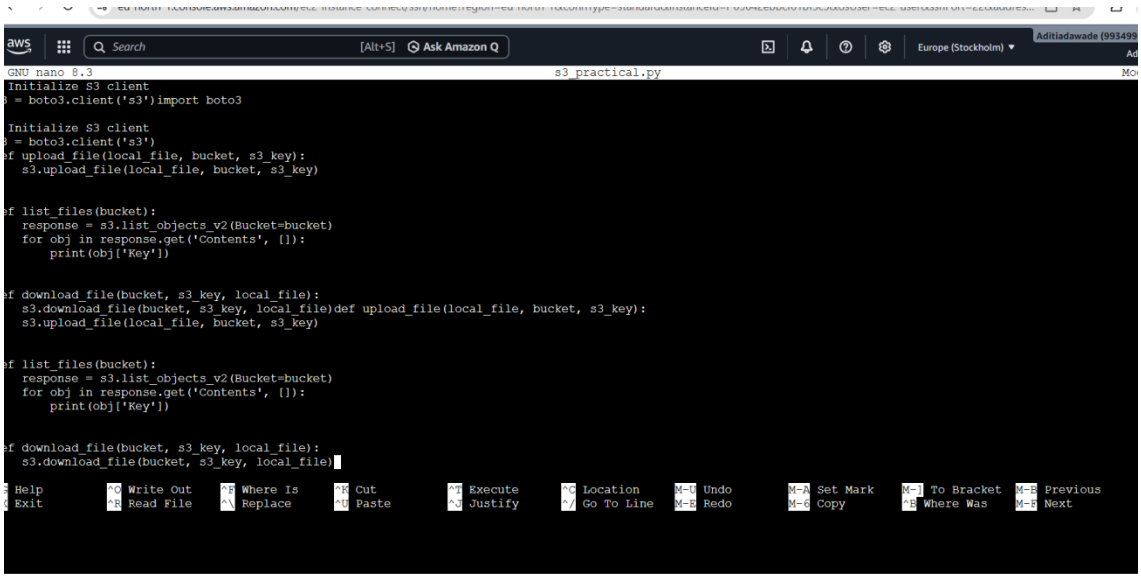


Step 4: Plan Your Script

Step 5: initialize the Boto3 S3 client



Step 6: Define Functions For each Operation

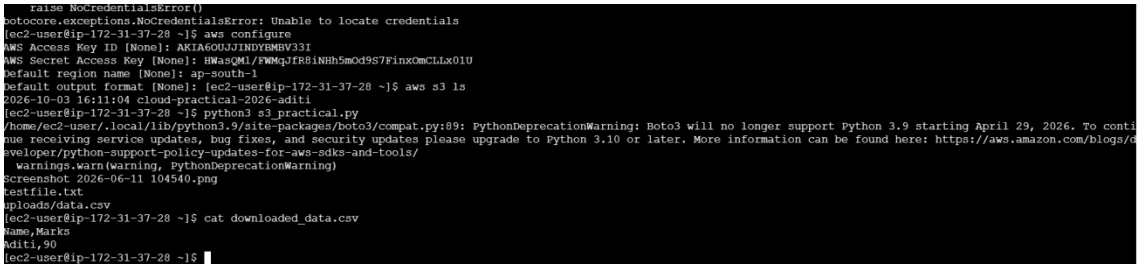


Step 7: Put it all together



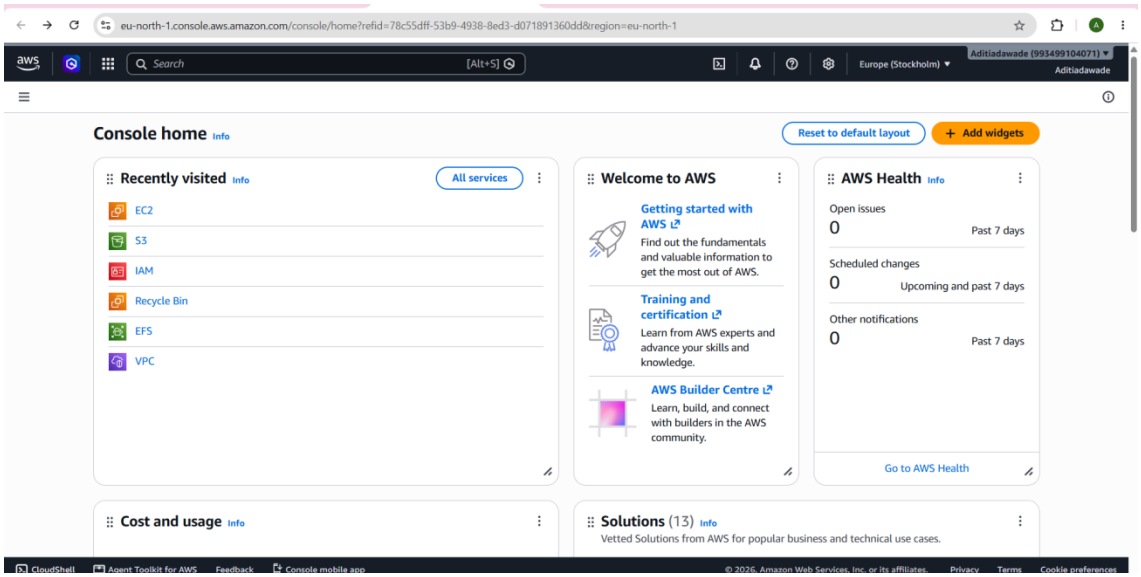
Step 8: Test and Debug

Step 9: Secure your Script

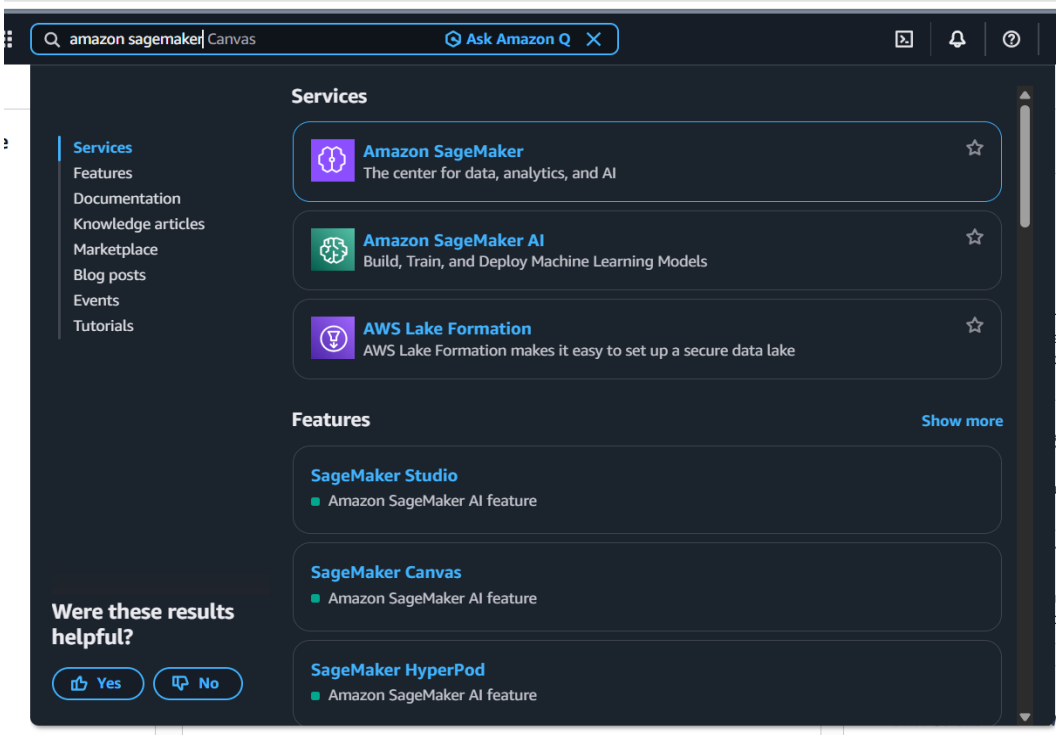


Practical no. 10

Step 1: Log in to the AWS Management console



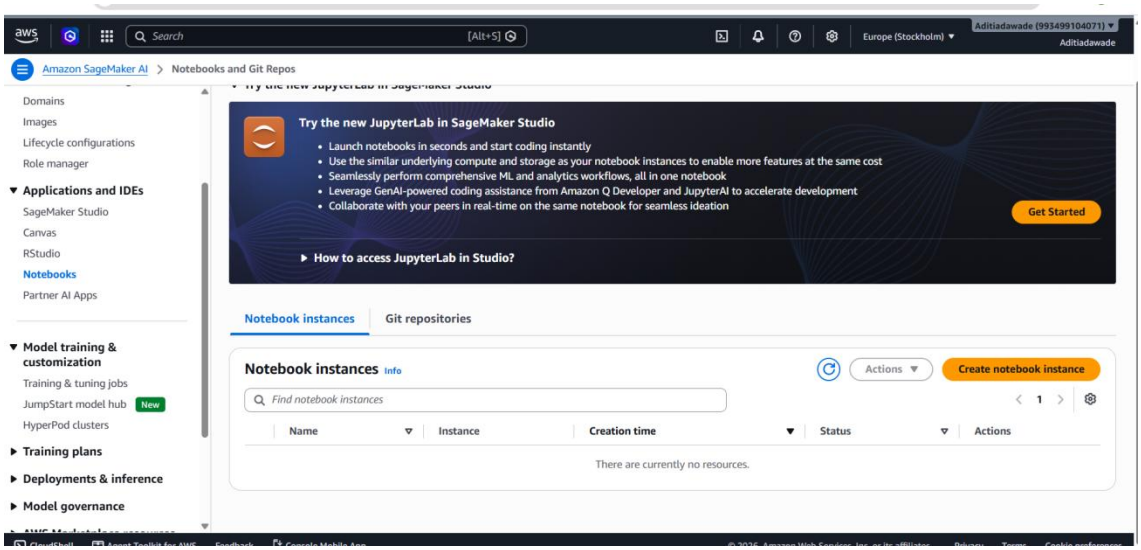
Step 2: Navigate to amazon SageMaker

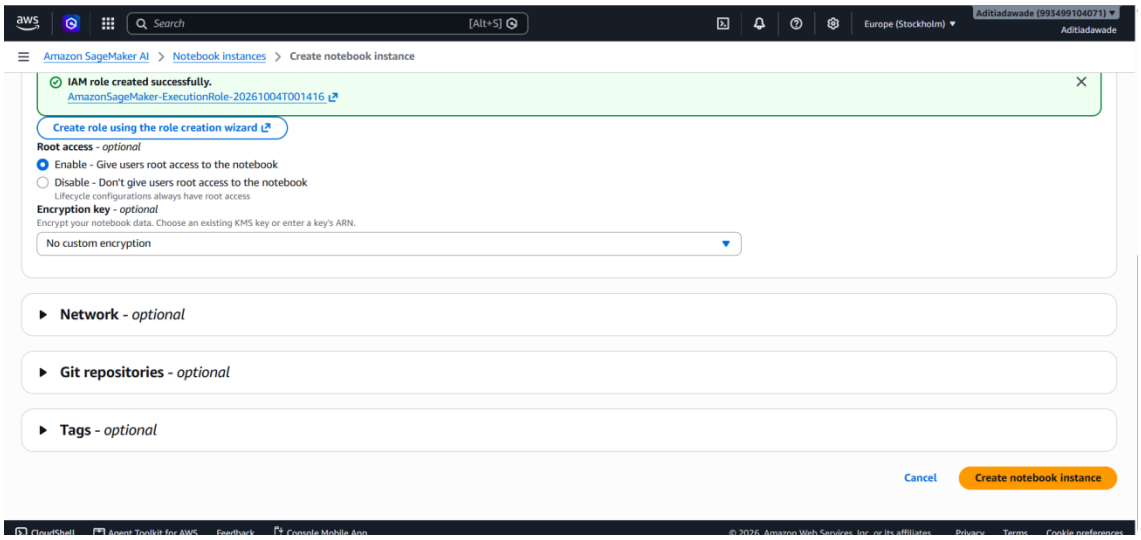


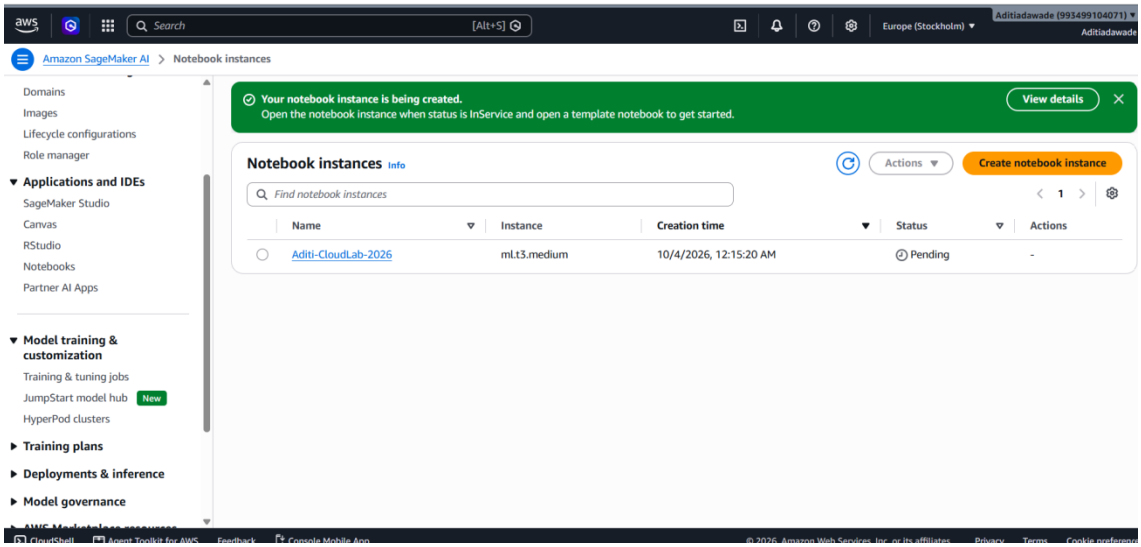
Step 3: Launch a notebook instance

Step 4: Configure the notebook

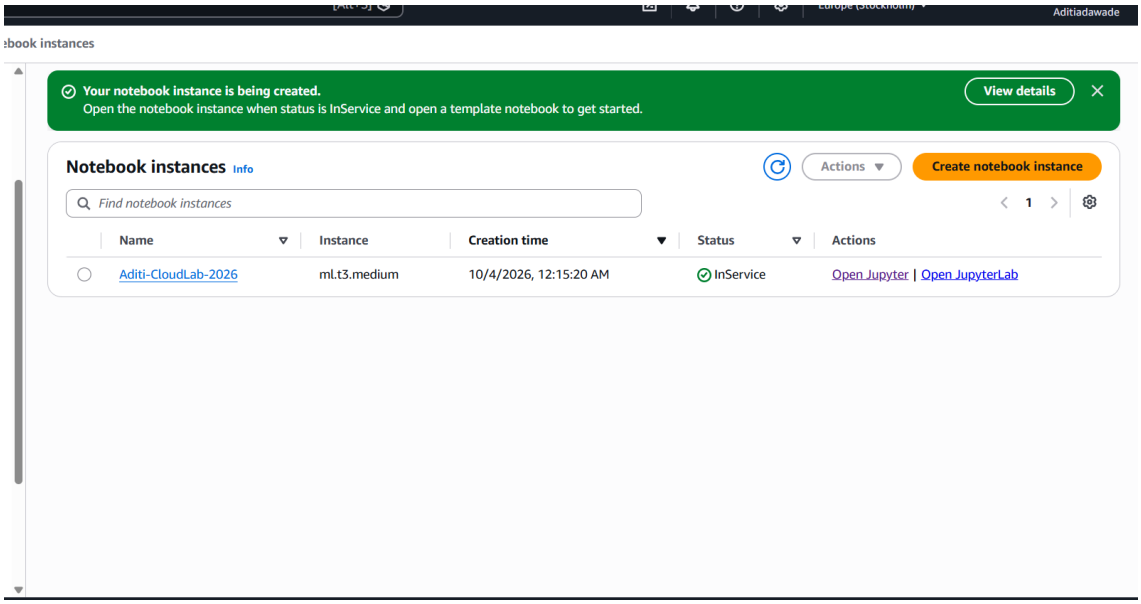
Step 5: Create the notebook



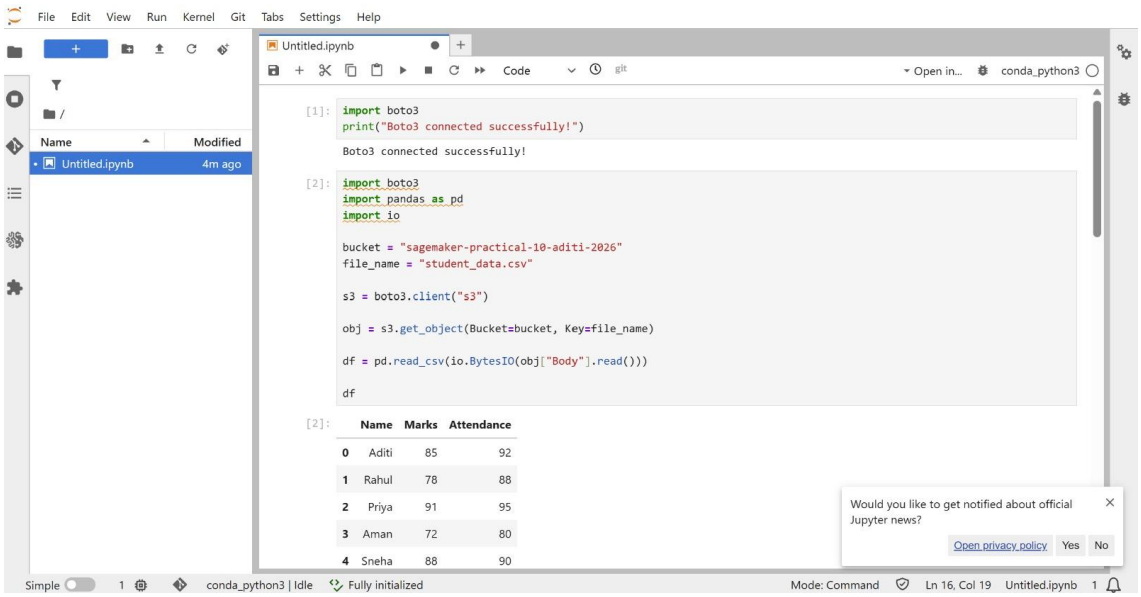


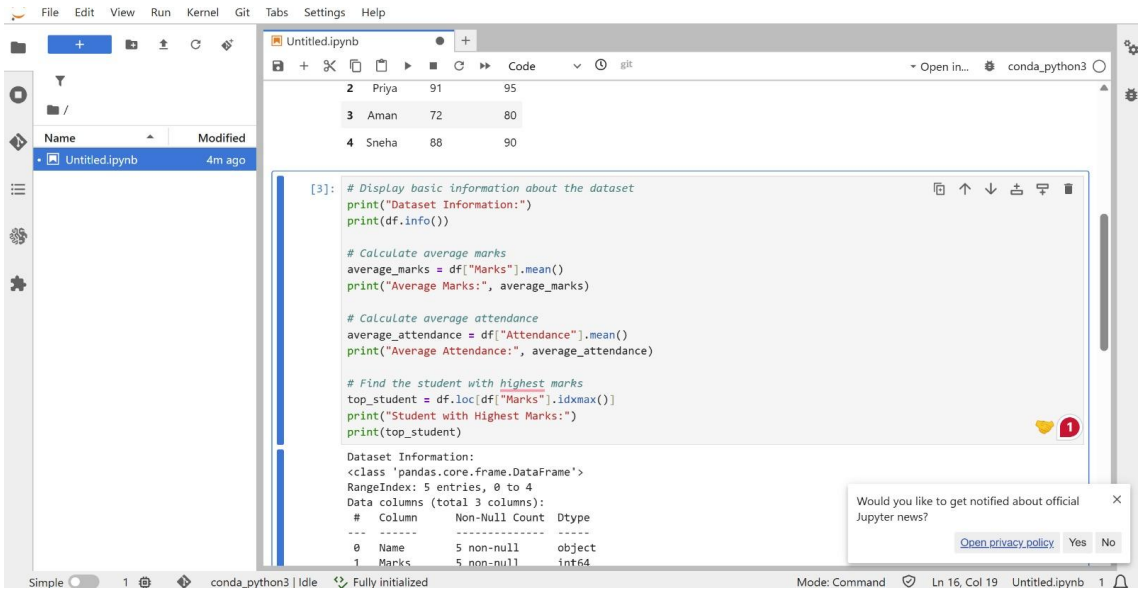


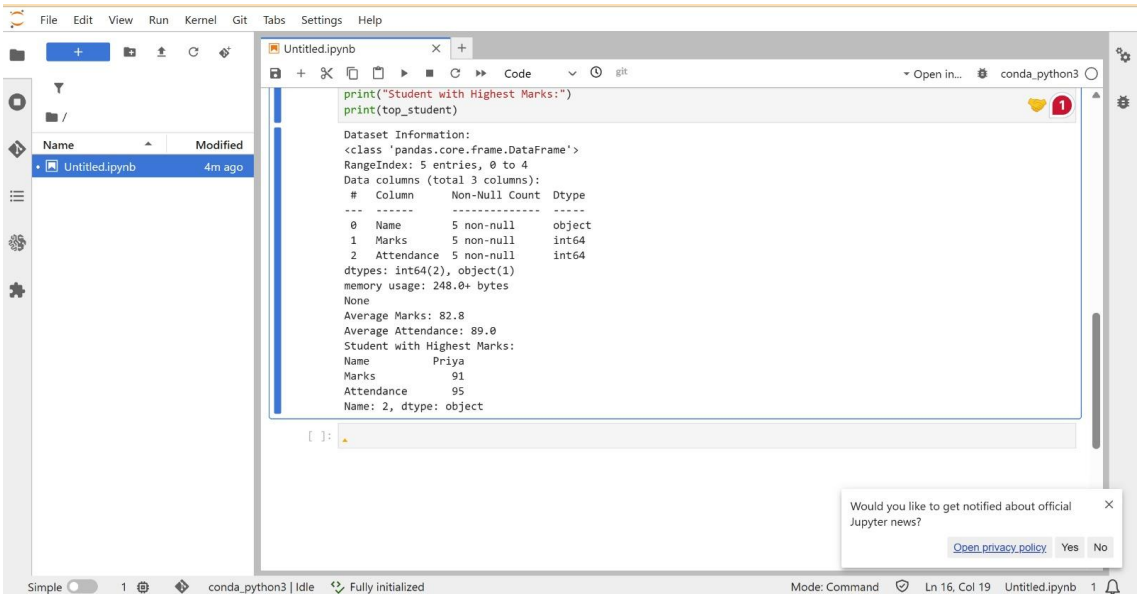
Step 6: Wait for the instance to Be ready



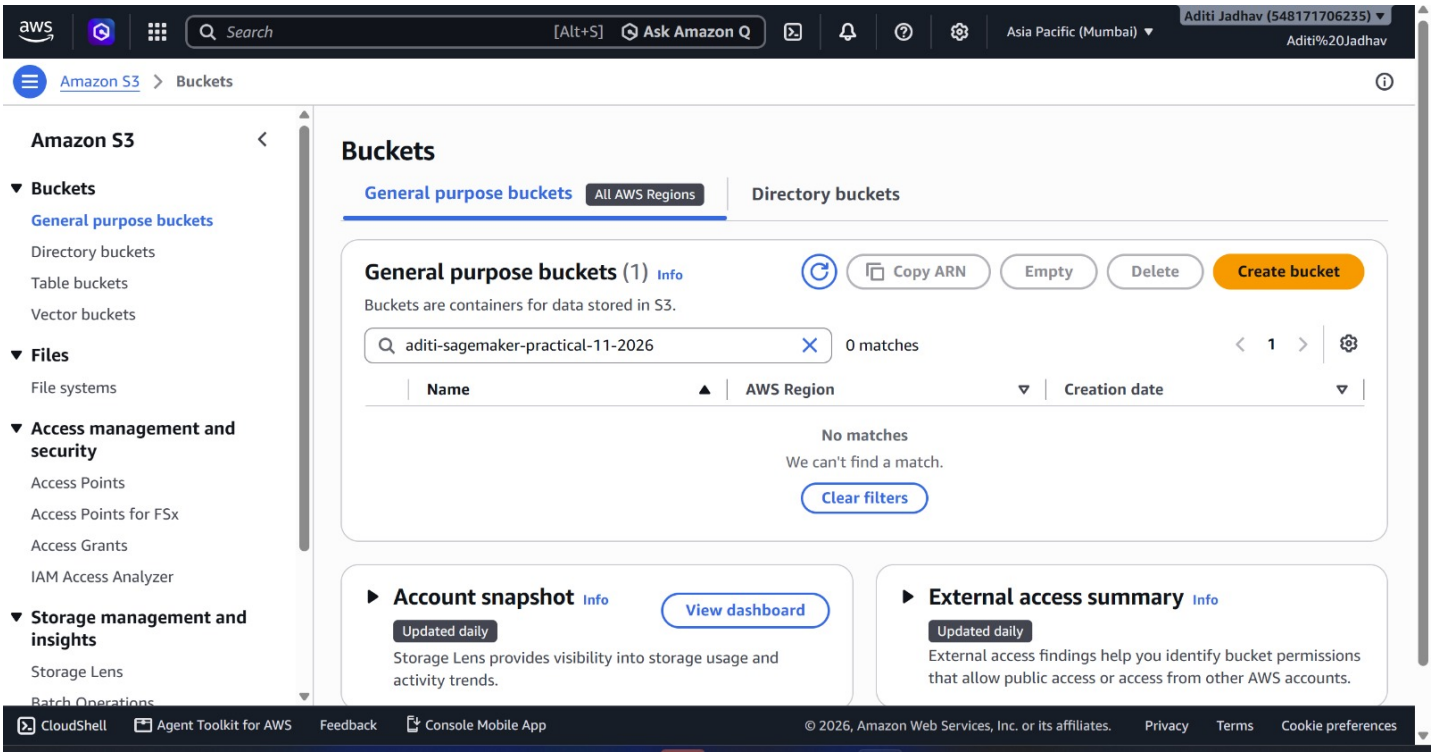
Step 7: Start using your notebook



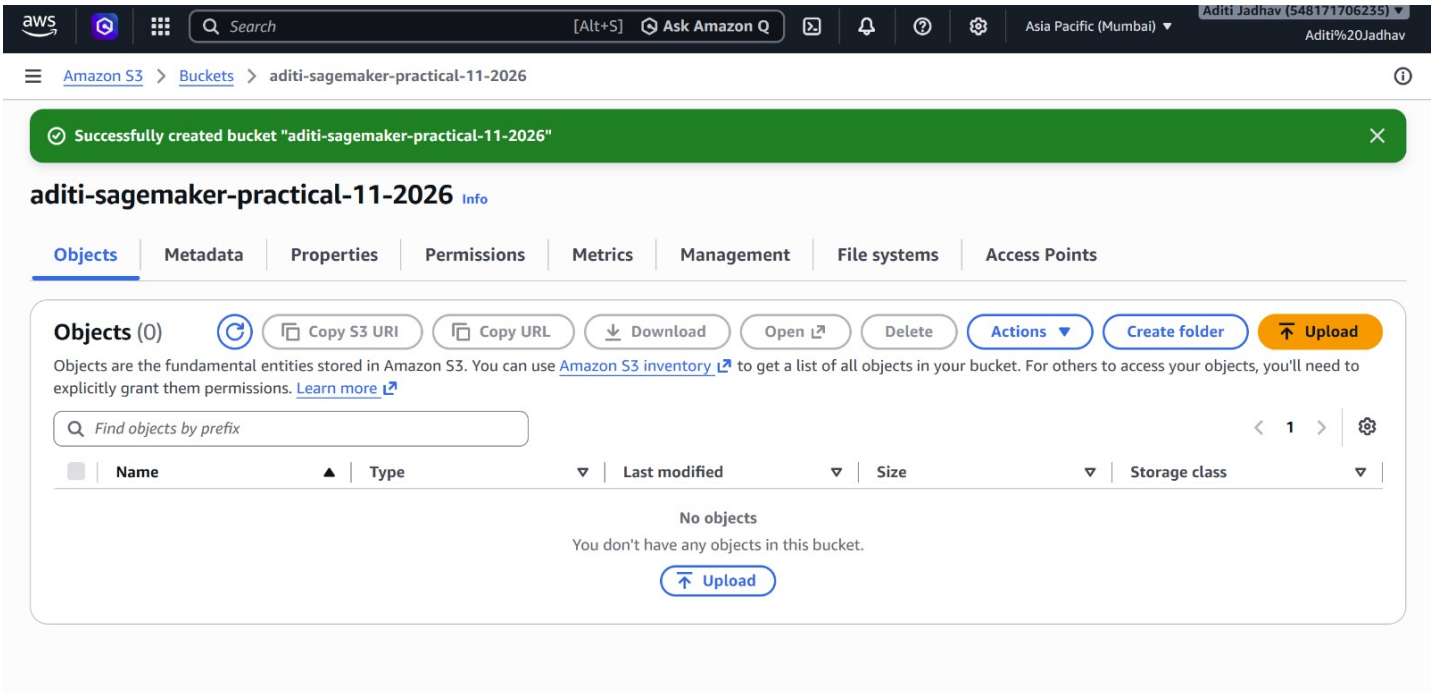



File Edit View Run Kernel Git Tabs Settings Help
Untitled.ipynb
Code
Open in...
conda_python3
Name
Modified
Untitled.ipynb
4m ago
print("Student with Highest Marks:")
print(top_student)
Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 3 columns):
Column Non-Null Count Dtype
0 Name 5 non-null object
1 Marks 5 non-null int64
2 Attendance 5 non-null int64
dtypes: int64(2), object(1)
memory usage: 248.0+ bytes
None
Average Marks: 82.8
Average Attendance: 89.0
Student with Highest Marks:
Name Priya
Marks 91
Attendance 95
Name: 2, dtype: object
Would you like to get notified about official Jupyter news?
Open privacy policy
Yes
No
Simple
conda_python3 | Idle
Fully initialized
Mode: Command
Ln 16, Col 19
Untitled.ipynb

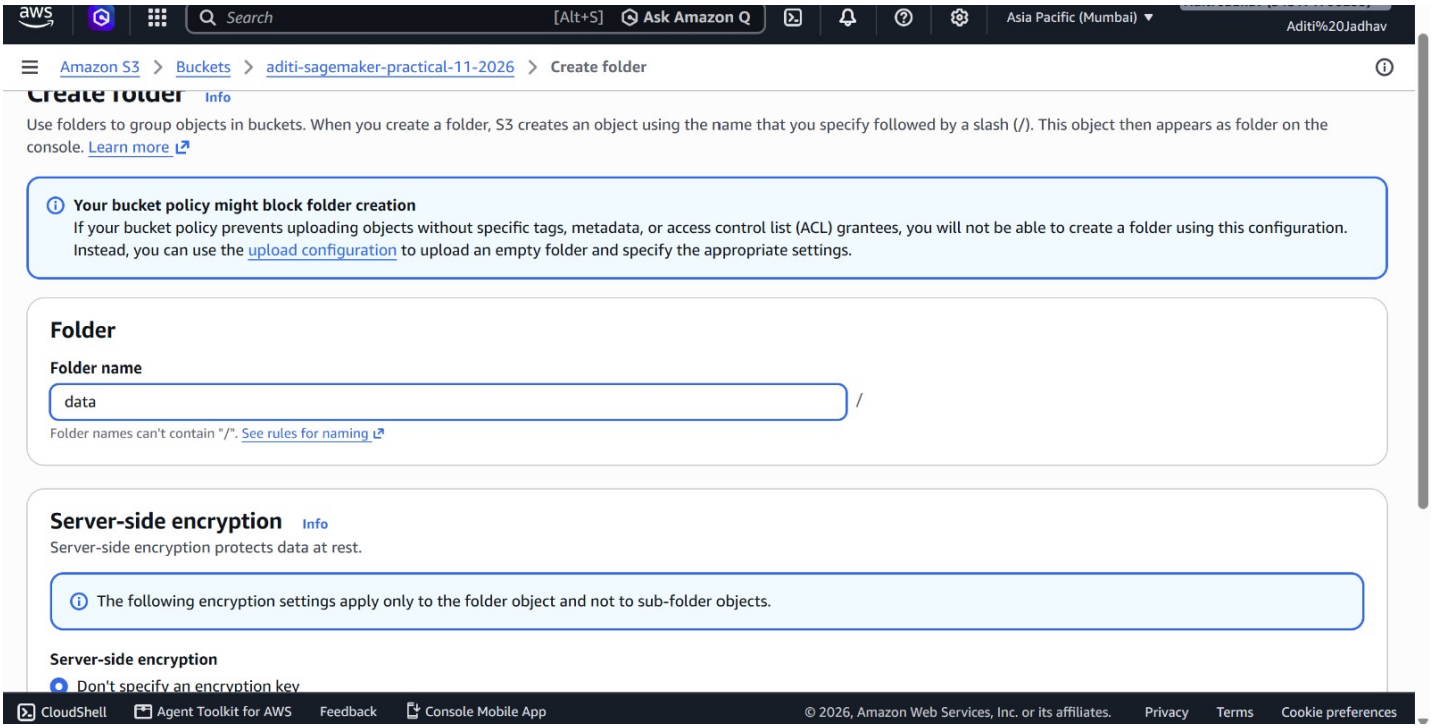
Screenshot 1: S3 bucket search: no matching bucket shown



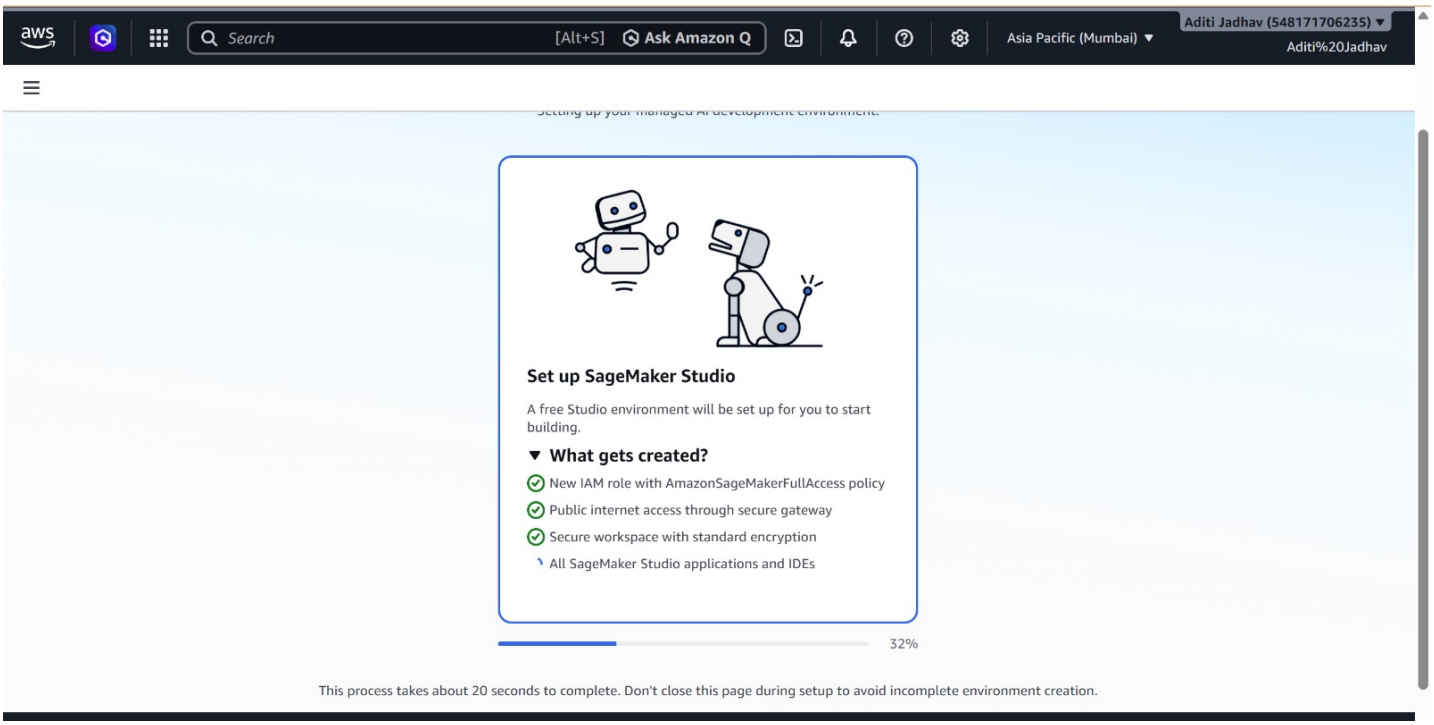
Screenshot 2: S3 bucket created successfully



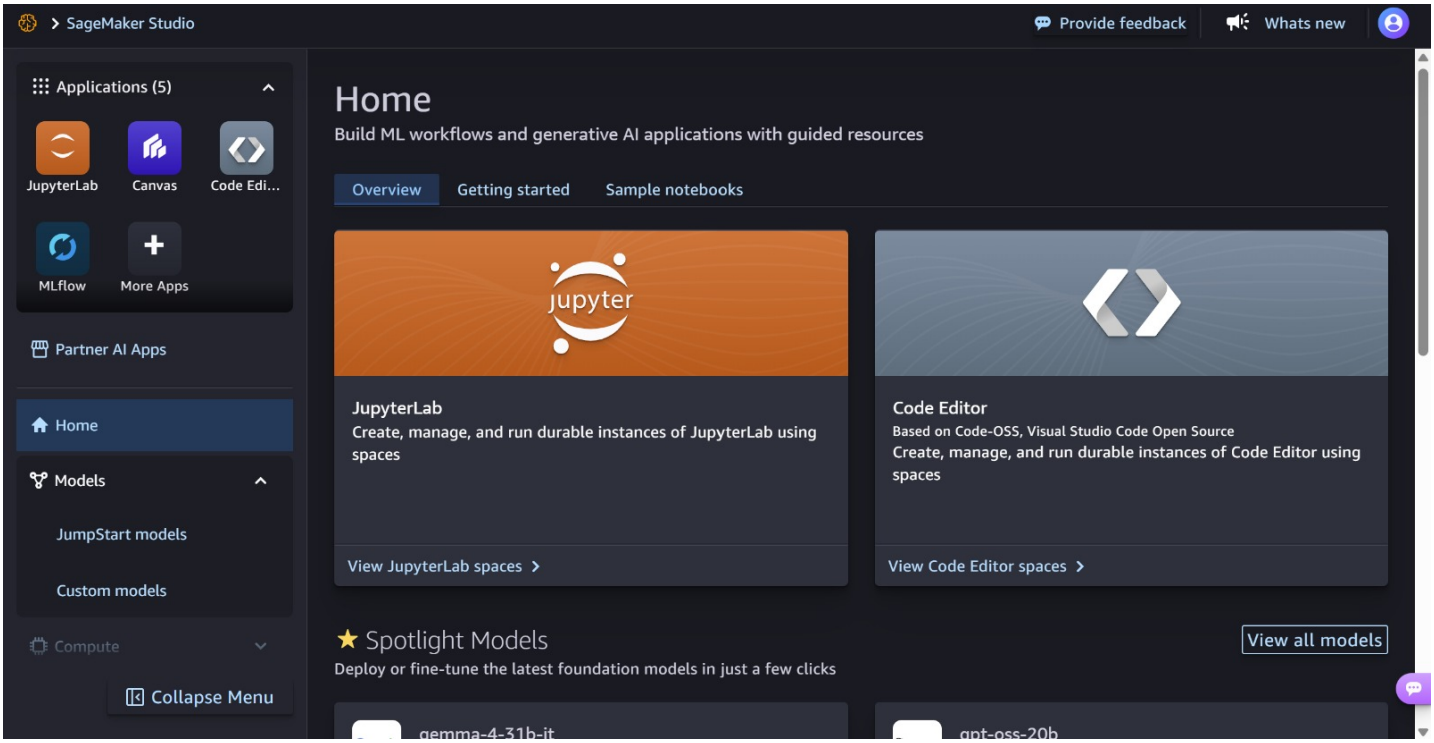
Screenshot 3: Create the data folder in the S3 bucket



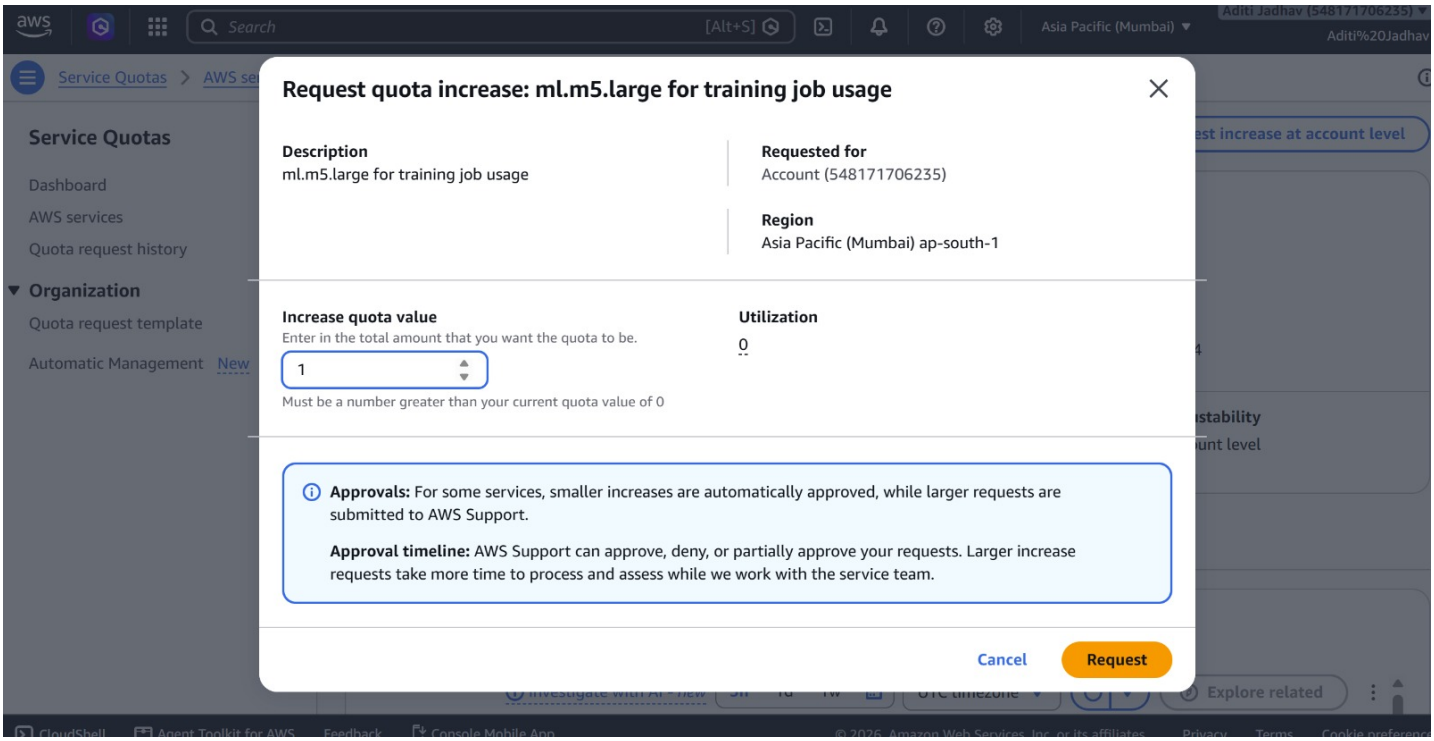
Screenshot 4: Set up SageMaker Studio



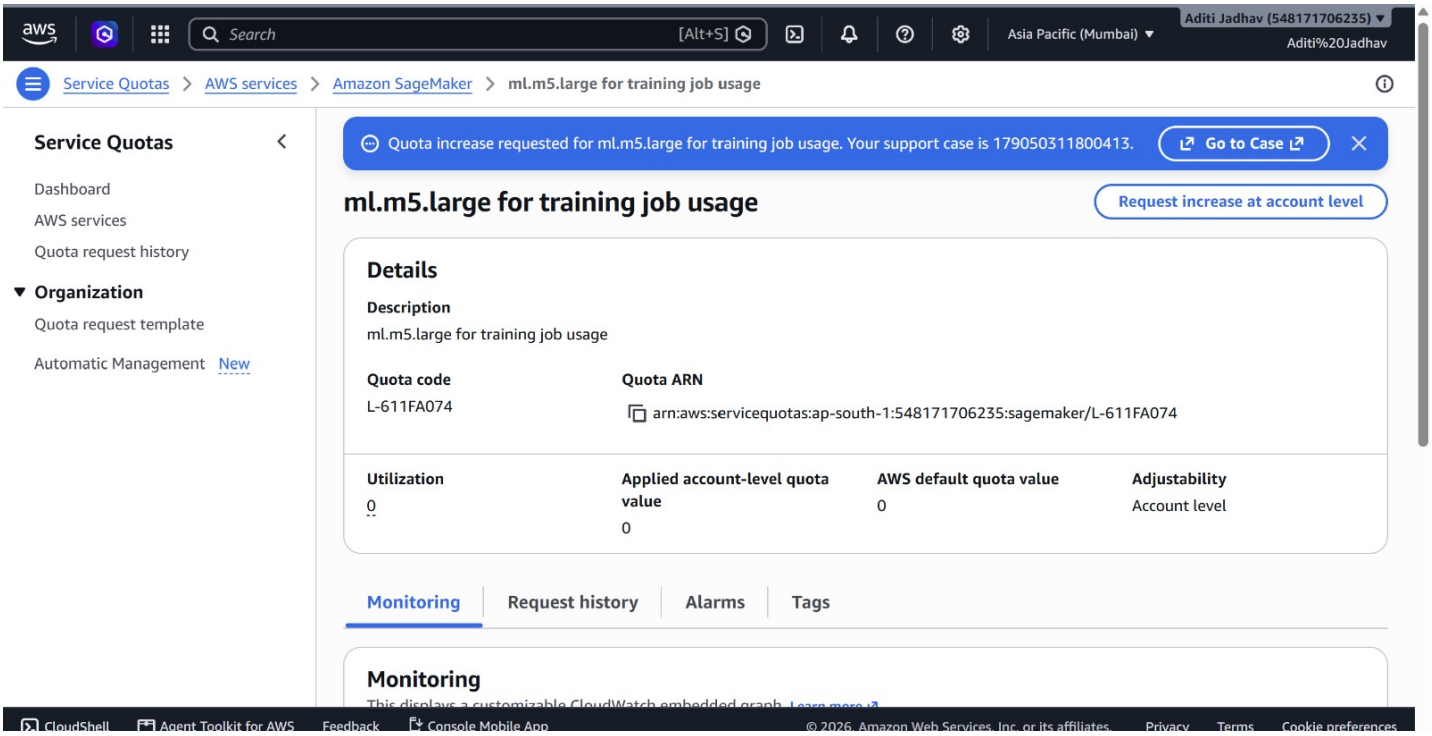
Screenshot 5: SageMaker Studio home



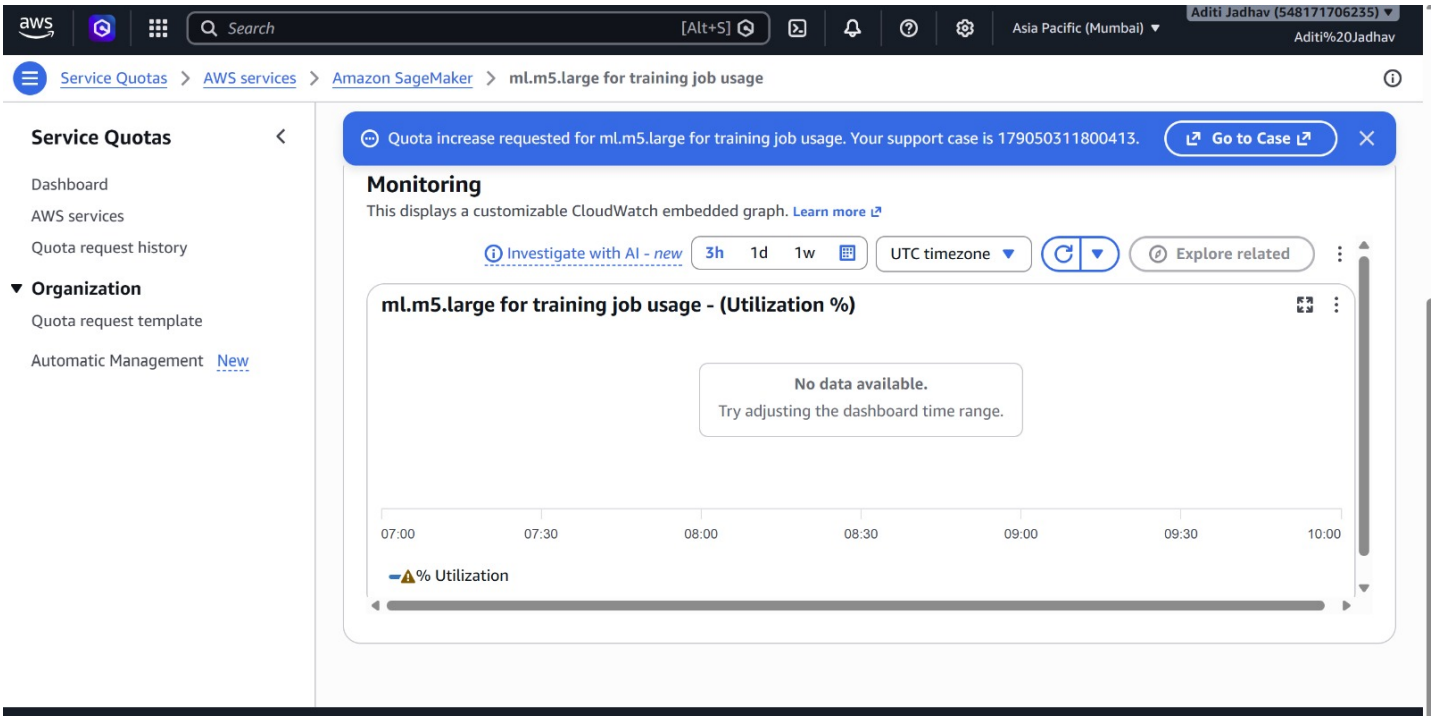
Screenshot 6: Request an ml.m5.large training-job quota increase



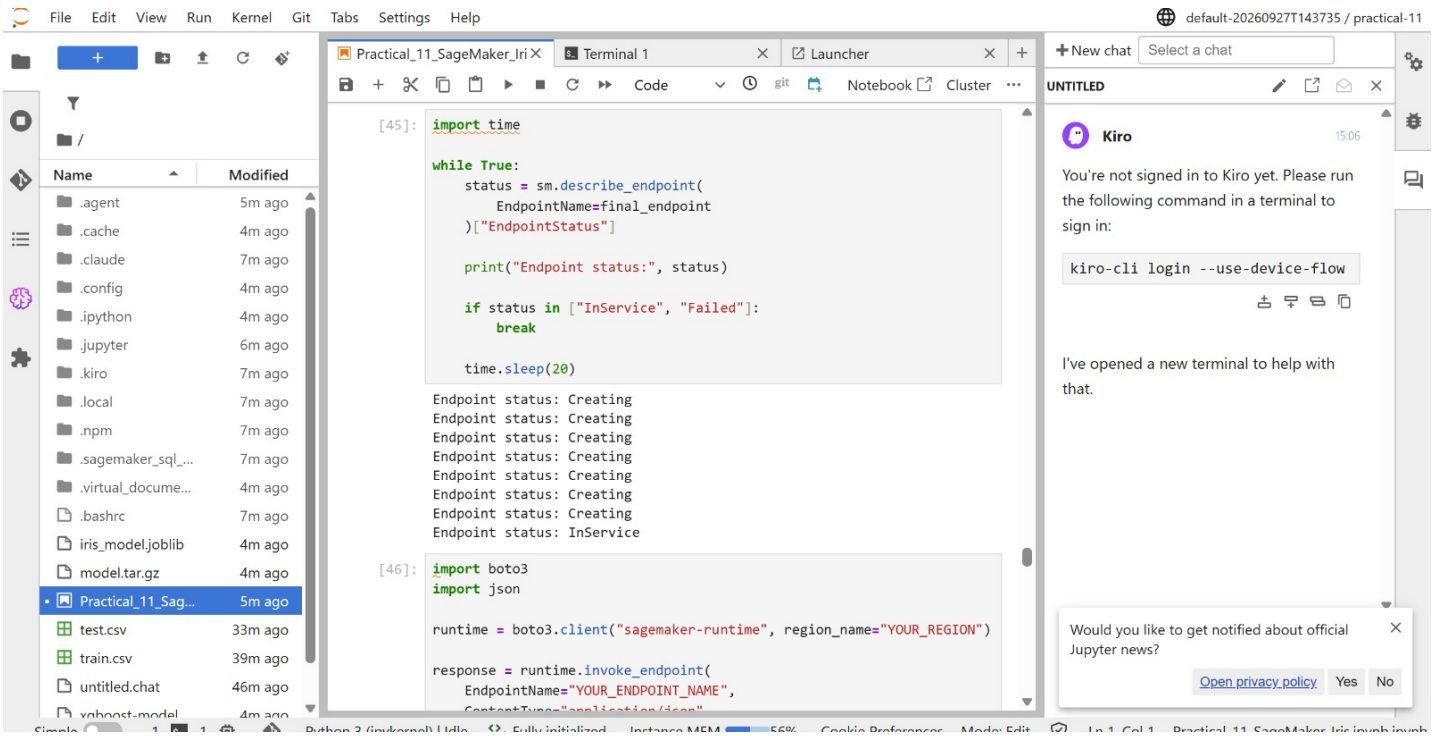
Screenshot 7: Training-job quota request confirmation



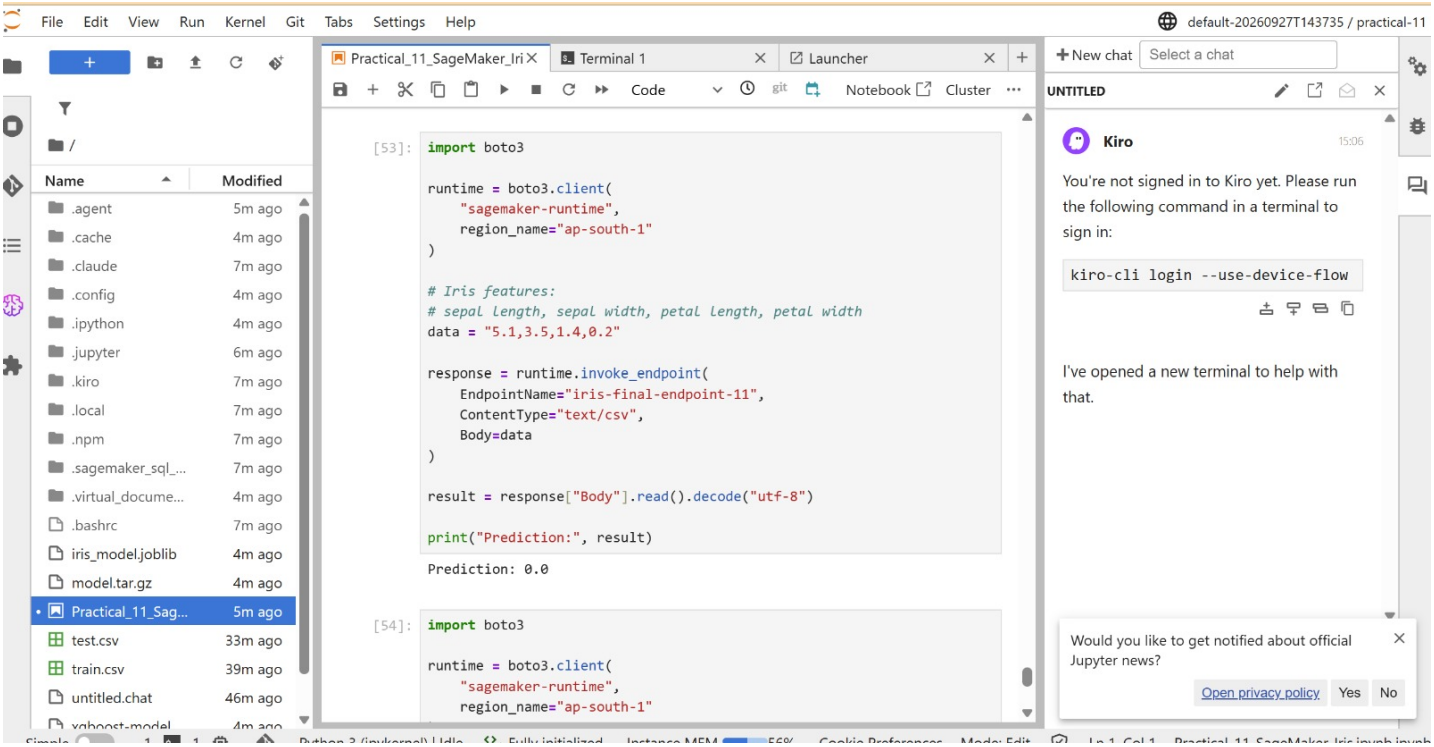
Screenshot 8: Training-job quota monitoring



Screenshot 9: Monitor the SageMaker endpoint creation status



Screenshot 10: Invoke the endpoint and view the prediction



Screenshot 11: Run another endpoint prediction

